

# Operating System: Introduction

---

Sang Ho Choi ([shchoi@kw.ac.kr](mailto:shchoi@kw.ac.kr))

School of Computer & Information Engineering  
KwangWoon University

# What happens when a program runs?

- A running program executes **instructions**
  1. The processor **fetches** an instruction from memory
  2. **Decode**: Figure out which instruction this is
  3. **Execute**: i.e., add two numbers, access memory, check a condition, jump to function, and so forth
  4. The processor moves on to the **next instruction** and so on

PC (program counter)

# Operating System (OS)

OS의 역할

- Responsible for
  - Making it easy to run programs
  - Allowing programs to share memory
  - Enabling programs to interact with devices

OS is in charge of making sure the system operates  
**correctly** and **efficiently**

# Virtualization

- The OS takes a **physical resource** and transforms it into a **virtual form** of itself
  - Physical resource: Processor, Memory, Disk ...
  - The virtual form is more general, powerful and easy-to-use
  - Sometimes, we refer to the OS as a **virtual machine**

현장의 자원을 가상화된 형태로 만들어서  
쉽게 사용하도록 만드는 것.

# System call

- System call allows user to tell the OS what to do  
    OS 가 어떤 일을 하도록 만드는 함수
    - The OS provides some interface (APIs, standard library)
    - A typical OS exports a few hundred system calls
- Run programs  
➤ Access memory  
➤ Access devices

# The OS is a resource manager

- The OS manage resources such as *CPU, memory* and *disk*      OS는 리소스 관리하고 조정하기도 함.
- The OS allows
  - Many programs to run → Sharing the CPU
  - Many programs to *concurrently* access their own instructions and data → Sharing memory
  - Many programs to access devices → Sharing disks

다수의 프로그램이 CPU, memory, disk를 이용해  
나누어에 따른 알고리즘이 포함

# Virtualizing the CPU

- The system has a very large number of virtual CPUs
  - Turning a single CPU into a seemingly infinite number of CPUs
  - Allowing many programs to seemingly run at once  
→ **Virtualizing the CPU**

ex) CPU는 4개일 때      다수의 프로그램이 이를 다 쓰려야 한다고  
     ttl 예상이 되어 있는 것처럼 보이게 함 (가상화)

# Virtualizing the CPU (Cont.)

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sys/time.h>
4 #include <assert.h>
5 #include "common.h"
6
7 int
8 main(int argc, char *argv[])
9 {
10     if (argc != 2) {
11         fprintf(stderr, "usage: cpu <string>\n");
12         exit(1);
13     }
14     char *str = argv[1];
15     while (1) {
16         Spin(1); // Repeatedly checks the time and
17                     returns once it has run for a second
18         printf("%s\n", str);
19     }
20 }
```

**Simple Example(cpu.c): Code That Loops and Prints**

# Virtualizing the CPU (Cont.)

- Execution result 1

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
^C
prompt>
```

Run forever; Only by pressing “Control-c” can we halt the program

# Virtualizing the CPU (Cont.)

- Execution result 2

```
prompt> ./cpu A &; ./cpu B &; ./cpu C &; ./cpu D &
[1] 7353
[2] 7354
[3] 7355
[4] 7356
A
B
D
C
A
B
D
C
A
C
B
D
...
```

동시에 실행되는  
것처럼 보이는  
효과.

Even though we have only one processor, all four of  
programs seem to be running at the same time!



# Virtualizing Memory

---

- The physical memory is *an array of bytes*
- A program keeps all of its data structures in memory
  - Read memory (load):
    - Specify an address to be able to access the data
  - Write memory (store):
    - Specify the data to be written to the given address

# Virtualizing Memory (Cont.)

- A program that Accesses Memory (mem.c)

```
1  #include <unistd.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include "common.h"
5
6  int
7  main(int argc, char *argv[])
8  {
9      int *p = malloc(sizeof(int)); // a1: allocate some
10     memory
11     assert(p != NULL);
12     printf("(%d) address of p: %08x\n",
13            getpid(), (unsigned) p); // a2: print out the
14     address of the memory
15     *p = 0; // a3: put zero into the first slot of the memory
16     while (1) {
17         Spin(1);
18         *p = *p + 1;
19         printf("(%d) p: %d\n", getpid(), *p); // a4
20     }
21     return 0;
22 }
```

# Virtualizing Memory (Cont.)

- The output of the program mem.c

```
prompt> ./mem
(2134) address of p: 00200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- The newly allocated memory is at address 00200000
- It updates the value and prints out the result

# Virtualizing Memory (Cont.)

- Running mem.c multiple times

```
prompt> ./mem &; ./mem &
[1] 24113
[2] 24114
(24113) memory address of p: 00200000
(24114) memory address of p: 00200000
(24113) p: 1
(24114) p: 1
(24114) p: 2
(24113) p: 2
(24113) p: 3
(24114) p: 3
...
...
```

- It is as if each running program has its own private memory
  - Each running program has allocated memory at the same address
  - Each seems to be updating the value at 00200000 independently

# Virtualizing Memory (Cont.)

- Each process accesses its own private virtual address space
  - The OS maps address space onto the physical memory
  - A memory reference within one running program does not affect the address space of other processes 32비트 주소공간이 영향을 X.
  - Physical memory is a shared resource, managed by the OS

# The problem of Concurrency

---

- The OS is juggling **many things at once**, first running one process, then another, and so forth
- Modern **multi-threaded programs** also exhibit the concurrency problem

# Concurrency Example

- A Multi-threaded Program (thread.c)

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include "common.h"
4
5  volatile int counter = 0;
6  int loops;
7
8  void *worker(void *arg) {
9      int i;
10     for (i = 0; i < loops; i++) {
11         counter++;
12     }
13     return NULL;
14 }
15 ...
```

# Concurrency Example (Cont.)

```
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
23         loops = atoi(argv[1]);
24         pthread_t p1, p2;
25         printf("Initial value : %d\n", counter);
26
27         Pthread_create(&p1, NULL, worker, NULL);
28         Pthread_create(&p2, NULL, worker, NULL);
29         Pthread_join(p1, NULL);
30         Pthread_join(p2, NULL);
31         printf("Final value : %d\n", counter);
32         return 0;
33     }
```

- The main program creates two threads
  - Thread: a function running within the same memory space. Each thread start running in a routine called `worker()`
  - `worker()`: increments a counter

# Concurrency Example (Cont.)

- loops determines how many times each of the two workers will increment the shared counter in a loop
  - loops: 1000

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- loops: 100000

200000 이 안나온다.

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

# Why is this happening?

- Increment a shared counter → take three instructions
  1. Load the value of the counter from memory into register
  2. Increment it
  3. Store it back into memory
- These three instructions do not execute **atomically** → Problem of concurrency happen

원자화 불가

제거와 삽입이 동시에 실행되는 경우 다른 프로세스의 영향에 의해 결과가 달라질 수 있다.

race condition

# Persistence

회원님

- Devices such as DRAM store values in a volatile
- *Hardware* and *software* are needed to store data **persistently**
  - **Hardware:** I/O device such as a hard drive, solid-state drives (SSDs)
  - **Software:**
    - File system manages the disk
    - File system is responsible for storing any files the user creates

# Persistence (Cont.)

- Create a file (/tmp/file) that contains the string "hello world"

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <assert.h>
4 #include <fcntl.h>
5 #include <sys/types.h>
6
7 int
8 main(int argc, char *argv[])
9 {
10     int fd = open("/tmp/file", O_WRONLY | O_CREAT
11                  | O_TRUNC, S_IRWXU);
12     assert(fd > -1);
13     int rc = write(fd, "hello world\n", 13);
14     assert(rc == 13);
15     close(fd);
16     return 0;
17 }
```

open(), write(), and close() system calls are routed to the part of OS called the file system, which handles the requests.

데이터가 소멸되지 않고  
기록으로 존재할 수 있게  
끔하는 OS의  
too)이 존재한다!

# Persistence (Cont.)

---

- What OS does in order to write to disk?
    - Figure out where on disk this new data will reside
    - Issue I/O requests to the underlying storage device
  - File system handles system crashes during write
    - Journaling or copy-on-write
    - Carefully ordering writes to disk
- 이스코어 2728점  
수록하는 기법.

# Design Goals

운영체계 설계 목표.

- Build up abstraction *사용하기 편리하게.*
  - Make the system **convenient and easy to use**
- Provide high performance *성능이 높아야 함.*
  - Minimize the overhead of the OS
  - OS must strive to provide virtualization without excessive overhead
- Protection between applications *프로그램 간 간섭을 X*
  - Isolation: Bad behavior of one does not harm other and the OS itself

# Design Goals (Cont.)

- High degree of reliability
  - The OS must also run non-stop
- Other issues
  - Energy-efficiency
  - Security
  - Mobility

신뢰성 ↑

제작 및 유지보수 편리

부가적인 추가 요소 .

중첩적인 부분은  
OS의 virtualization, concurrently, persistence.