

Operating System: Processes

Sang Ho Choi (shchoi@kw.ac.kr)

School of Computer & Information Engineering
KwangWoon University


Virtualization

A Dialogue on Virtualization

- Professor: And thus we reach the first of our three pieces on operating systems: virtualization
- Student: But what is virtualization, oh noble professor?
- Professor: Imagine we have a peach



A Dialogue on Virtualization (Cont'd)




A Dialogue on Virtualization (Cont'd)

- Student: Well, if I was sharing a peach with somebody else, I think I would notice
- Professor: Ah yes! Good point. But that is the thing with many eaters: **most of the time they are napping or doing something else**, and thus, you can snatch that peach away and give it to someone else for a while. And thus we create **the illusion of many virtual peaches, one peach for each person!**


CPU Virtualization

기상화를 하기 위한 단계 : context switching
▷ 기상 : scheduling


The Crux of the Problem:
How to Provide the Illusion of Many CPUs?



Running a Process




Running Multiple Processes




Interleaving Multiple Processes

process가 2개만 있는
경우에는 문제 X

하나가 2개로 나눠야 하는
경우에는 문제 X



이거는 scheduling을
진행하는 알고리즘을
가지는 걸 아닐까?
"OS"





Virtualizing the CPU

Time sharing of the CPU

(time slicing
time quantum)

사용시간은 너무 짧아
잡으면 전환의 번거
이 너무 많아 안 좋고
길면 하나가 늑장함.
따라서 알고리즘이 필요.



OS creates the illusion
that each process has its
own CPU (and memory)

각각의 프로세스가 각각
CPU를 할당받는 것처럼
만들어준다.



광운대학교
KwangWoon University

A Process

A process is a **running program**

- Comprising of a process:
 - Memory (address space)
 - Code(Text)
 - Data
 - Stack
 - Heap
 - Registers
 - Program counter
 - Stack pointer

Process API

- These APIs are available on any modern OS
 - Create
 - Create a new process to run a program
 - Destroy
 - Halt a runaway process
 - Wait
 - Wait for a process to stop running
 - Miscellaneous Control *비정상적인 종료*.
 - Some kind of method to suspend a process and then resume it
 - Status *running인지* *block인지 등등...*
 - Get some status info about a process

Process Creation

1. Load a program code into memory, into the address space of the process
 - Programs initially **reside on disk** in **executable format**
 - OS perform the loading process **lazily**
 - Loading pieces of code or data only as they are needed during program execution *lazily*의 의미는 코드의 일부분만 메모리에 가져오고 필요할 때마다 가져온다는 것
2. The program's run-time **stack** is allocated *자체영역*.
 - Use the stack for *local variables*, *function parameters*, and *return address*
 - Initialize the stack with arguments → argc and the argv array of main () function

Process Creation (Cont.)

3. The program's heap is created

- Used for explicitly requested dynamically allocated data
- Program request such space by calling `malloc()` and free it by calling `free()`

프로그램의 힙이 생성된다.


4. The OS do some other initialization tasks

- input/output (I/O) setup
 - Each process by default has three open file descriptors
 - Standard input, output and error

5. Start the program running at the entry point, namely `main()`

- The OS *transfers control* of the CPU to the newly-created process


Loading: From Program To Process



Process States

- A process can be one of three states
 - Running CPU를 할당 받아서 실행중인 상태.
➤ A process is running on a processor
 - Ready 다른 프로세스가 사용 중이라 대기 중인 상태.
➤ A process is ready to run but for some reason the OS has chosen not to run it at this given moment
 - Blocked 특수한 일을 수행하기 위한 상태
➤ A process has performed some kind of operation
➤ When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor I/O 요청 시에는 CPU가 필요 없기 때문에.

Process State Transition



Tracing Process State: CPU Only

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

Tracing Process State: CPU and I/O

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Blocked	Running	Process ₀ initiates I/O
5	Blocked	Running	Process ₀ is blocked, so Process ₁ runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	—	
10	Running	—	Process ₀ now done

Data structures

- The OS has **some key data structures** that track various relevant pieces of information
 - Process list
 - Ready processes
 - Blocked processes
 - Current running process
 - Register context
- PCB (Process Control Block)
 - A C-structure that contains information **about each process**

Implementing Processes

- PCB (Process Control Block) or Process Descriptor
 - Contains all of the information about a process
 - CPU registers
 - PID, PPID, process group, priority, process state, signals
 - CPU scheduling information
 - Memory management information
 - Accounting information — ↗
해당 프로세스가 얼마나 사용되었는지
 - File management information
 - I/O status information
 - Credentials
 - task_struct in Linux
 - 3248 bytes as of Linux 3.2.0

Example) The xv6 kernel Proc Structure

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;      // Index pointer register
    int esp;      // Stack pointer register
    int ebx;      // Called the base register
    int ecx;      // Called the counter register
    int edx;      // Called the data register
    int esi;      // Source index register
    int edi;      // Destination index register
    int ebp;      // Stack base pointer register
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```

Example) The xv6 kernel Proc Structure (Cont.)

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                      // Start of process memory
    uint sz;                         // Size of process memory
    char *kstack;                    // Bottom of kernel stack
                                    // for this process
    enum proc_state state;          // Process state
    int pid;                         // Process ID
    struct proc *parent;            // Parent process
    void *chan;                      // If non-zero, sleeping on chan
    int killed;                      // If non-zero, have been killed
    struct file *ofile[NFILE];      // Open files
    struct inode *cwd;              // Current directory
    struct context context;          // Switch here to run process
    struct trapframe *tf;           // Trap frame for the
                                    // current interrupt
};
```

Context Switch

- When CPU switches to another process, the system must
 - save the state of the old process, and
 - load the saved state for the new process
- Context switch time is pure overhead
 - System does not do any useful work while switching
전환하는 시간이 유휴시간이 아니 time slicing이 가능하지 않으면 X.
- Context switch time depends on hardware
 - The register set is different

Context Switch (Cont.)

- CPU switch from process to process

