

Operating System: Limited Direct Execution

Sang Ho Choi (shchoi@kw.ac.kr)

School of Computer & Information Engineering
KwangWoon University

How to efficiently virtualize the CPU with control?

- The OS needs to share the physical CPU by time sharing
 - Issue *수동분기시에는 시스템적인 과도한 오버헤드를 어떻게 줄일 수 있는지?*
 - Performance: How can we implement virtualization without adding excessive overhead to the system?
 - Control: How can we run processes efficiently while retaining control over the CPU?*전트롤 분기시에는 그림을 더 효율적으로 사용하기 위한 방법은?*

Direct Execution

운영체제가 그냥 시스템에 접근하거나
직접 실행되는가 아래 OS는 그냥
라이브러리 수준의 일반하는 거.

- Just run the program directly on the CPU

OS	Program
<ol style="list-style-type: none">Create entry for process listAllocate memory for programLoad program into memorySet up stack with argc / argvClear registersExecute call main ()Free memory of processRemove from process list	<ol style="list-style-type: none">Run main ()Execute return from main ()

Without *limits* on running programs,
the OS wouldn't be in control of anything and
thus would be "**just a library**"

OS는 추가적인 운영체제에 대한 컨트롤을 할 줄 알아야 함.

Problem 1: Restricted Operation

- What if a process wishes to perform some kind of restricted operation such as ...
 - Issuing an I/O request to a disk
 - Gaining access to more system resources such as CPU or memory *할당 받은 자원 이상을 사용하지 않도록.*
- Solution: Using protected control transfer
 - **User mode**: Applications do not have full access to hardware resources.
 - **Kernel mode**: The OS has access to the full resources of the machine

권한을 나눠서 문제를 해결.

System Call

- Allow the kernel to **carefully expose** certain key pieces of functionality to user program, such as ...
 - Accessing the file system
 - Creating and destroying processes
 - Communicating with other processes
 - Allocating more memory

System Call (Cont.)

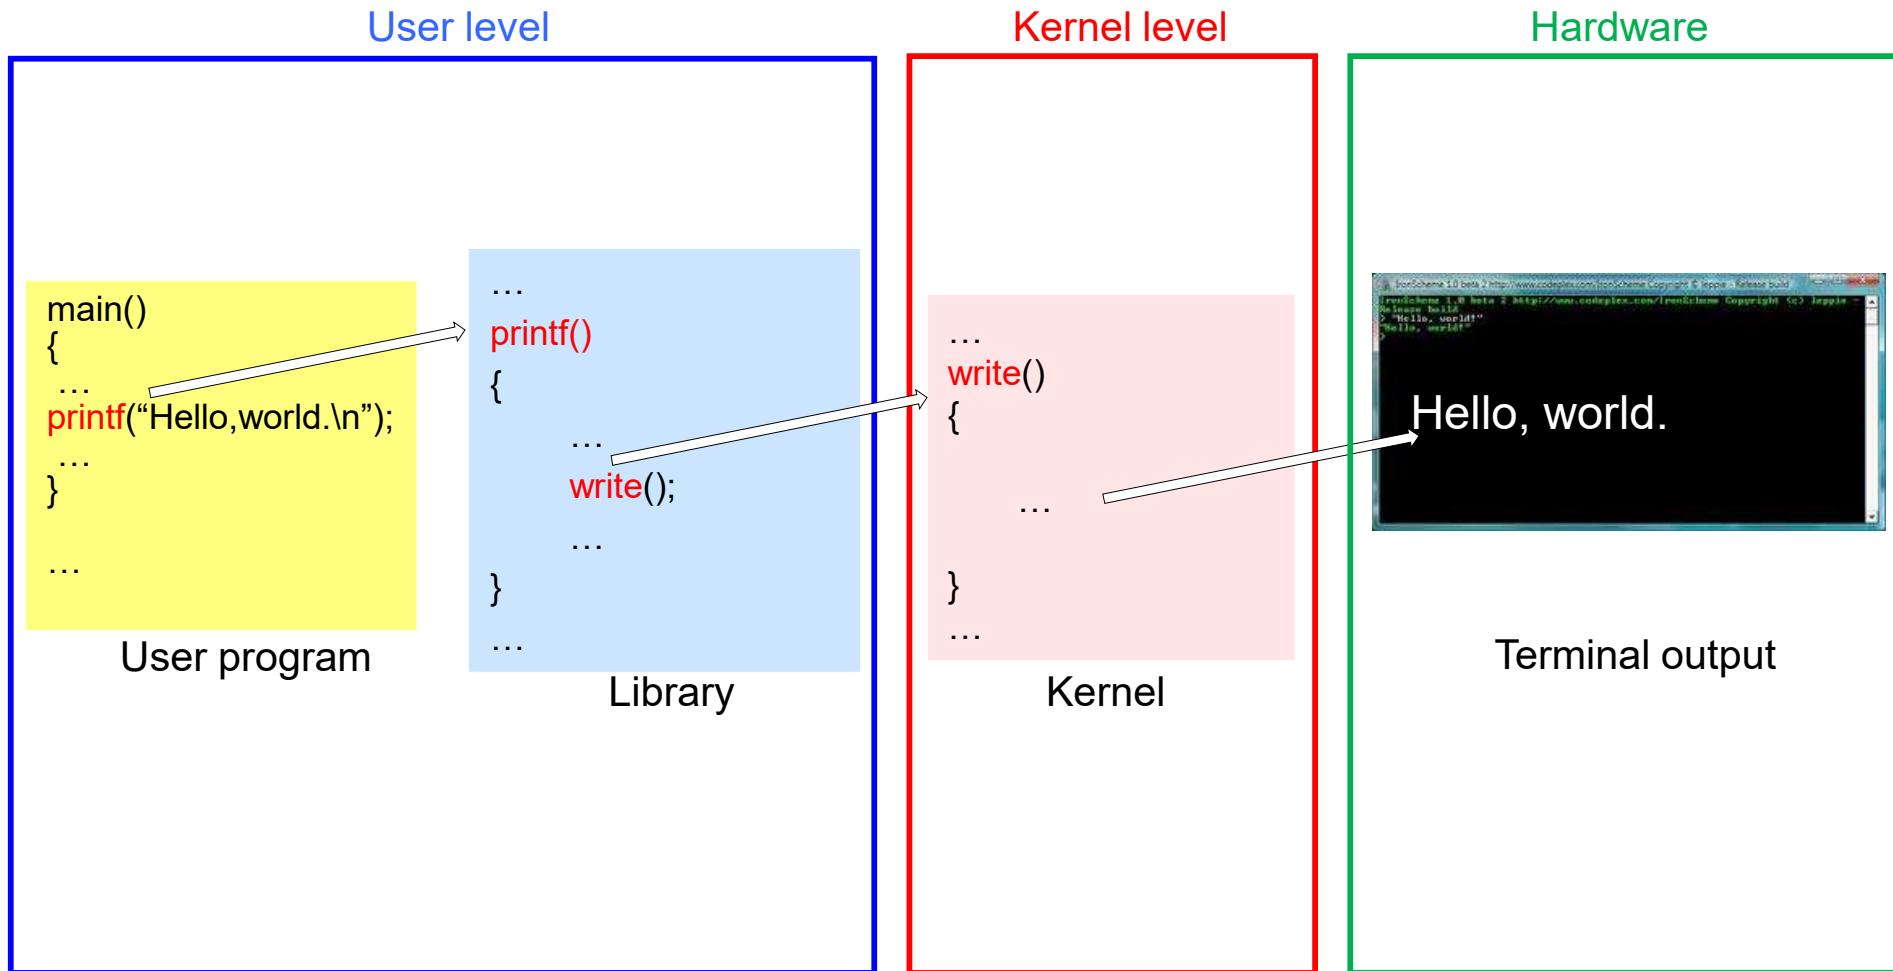
User mode → kernel mode
사용자모드 Trap.

- Trap instruction
 - Jump into the kernel
 - Raise the privilege level to kernel mode
- Return-from-trap instruction
 - Return into the calling user program
 - Reduce the privilege level back to user mode
- How does the trap know which code to run inside the OS?
 - trap table
 - also called interrupt descriptor table or interrupt vector table
 - trap handler
 - The code should be run when a program executes a trap instruction

System Call (Cont.)

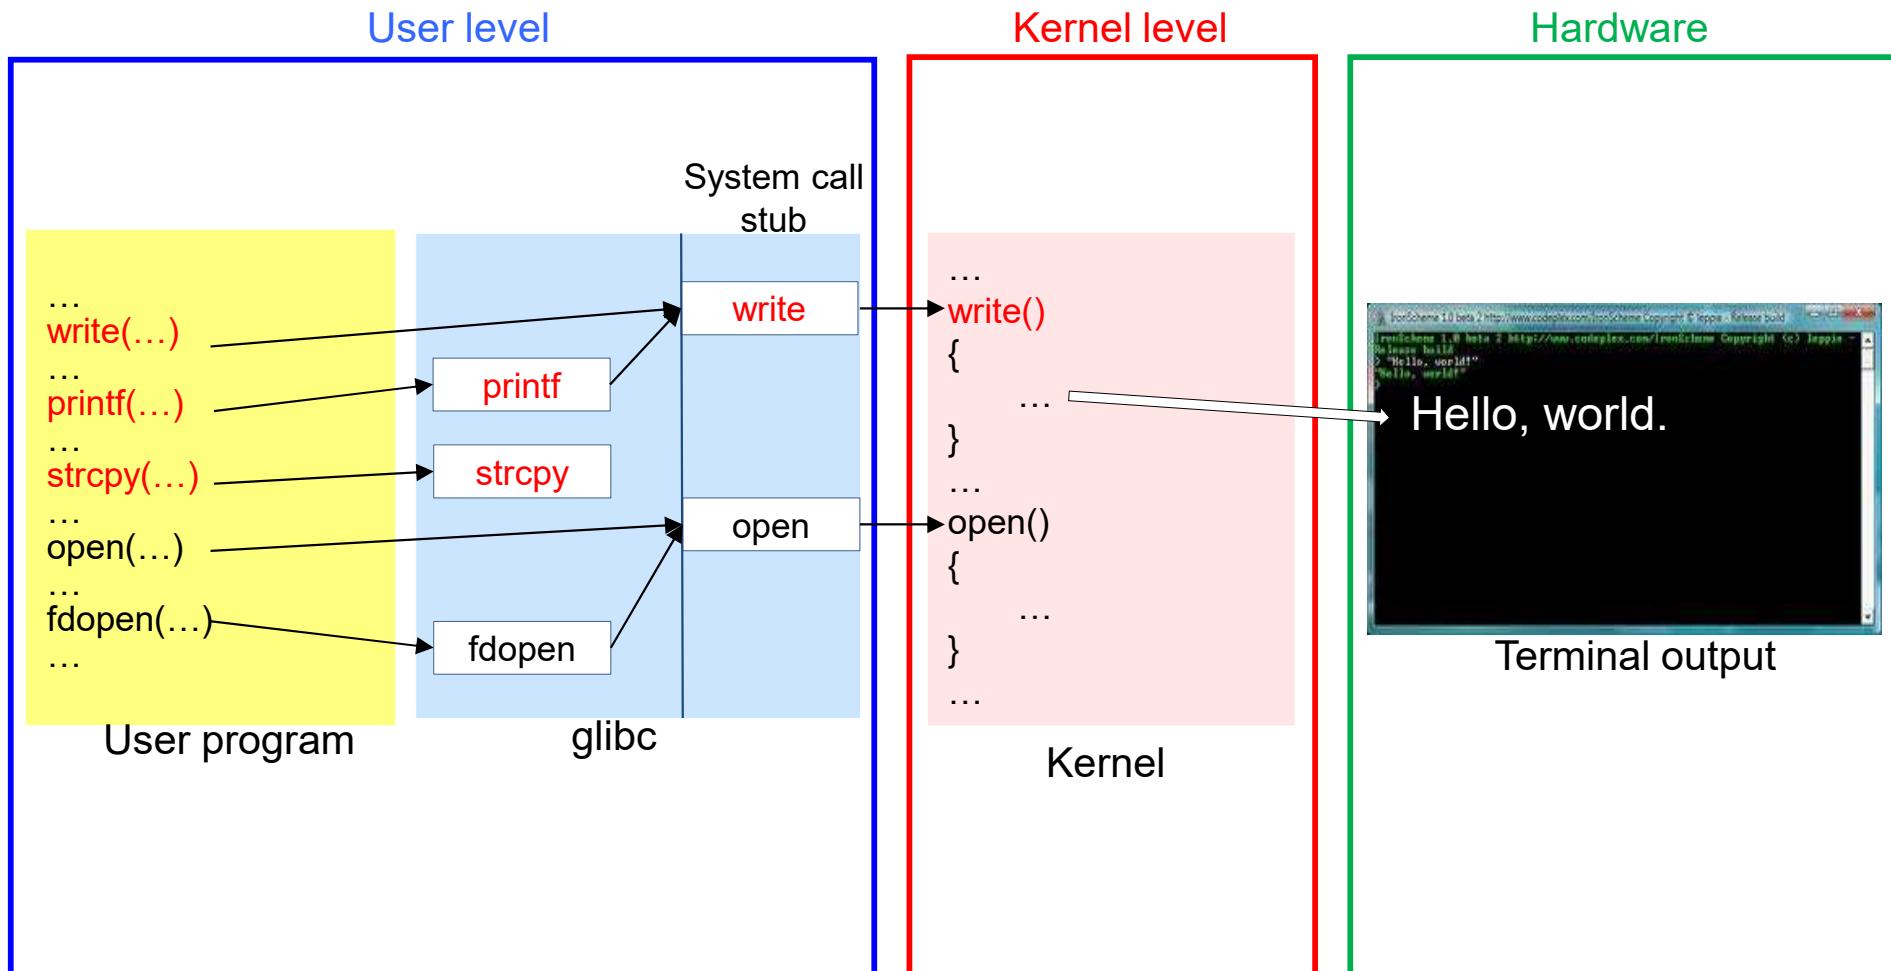
- System-call number
 - Assigned to each system call
 - The user code is thus responsible for placing the desired system-call number in a register

System Call Handling



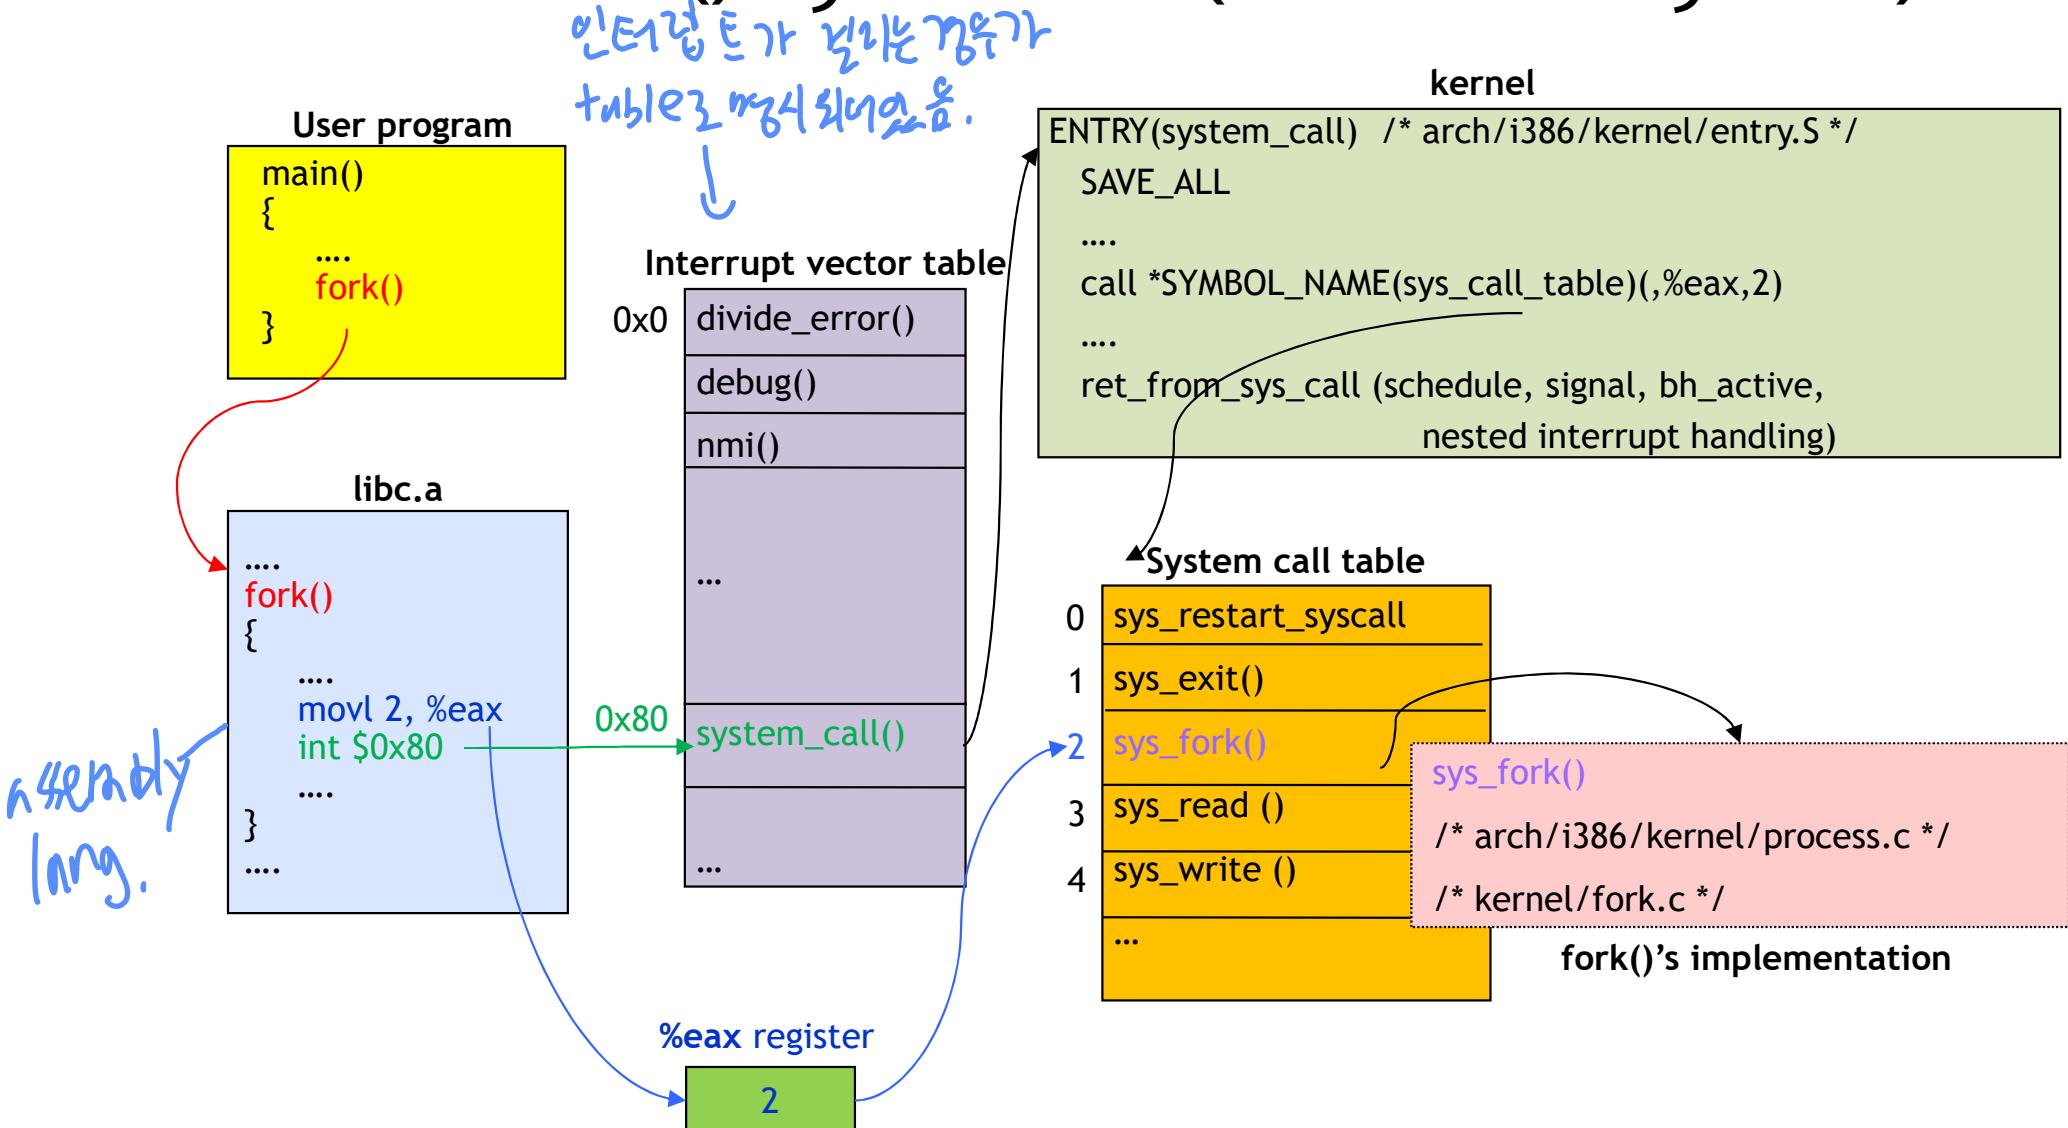
System Call Handling (Cont.)

user mode의 write가
kernel mode의 write를 탐.



System Call Handling (Cont.)

- Invoke fork() system call (i386-linux system)



Limited Direction Execution Protocol

OS @ boot
(kernel mode)

initialize trap table

Hardware

remember address of ...
syscall handler

OS @ run
(kernel mode)

Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC
return-from -trap

Hardware

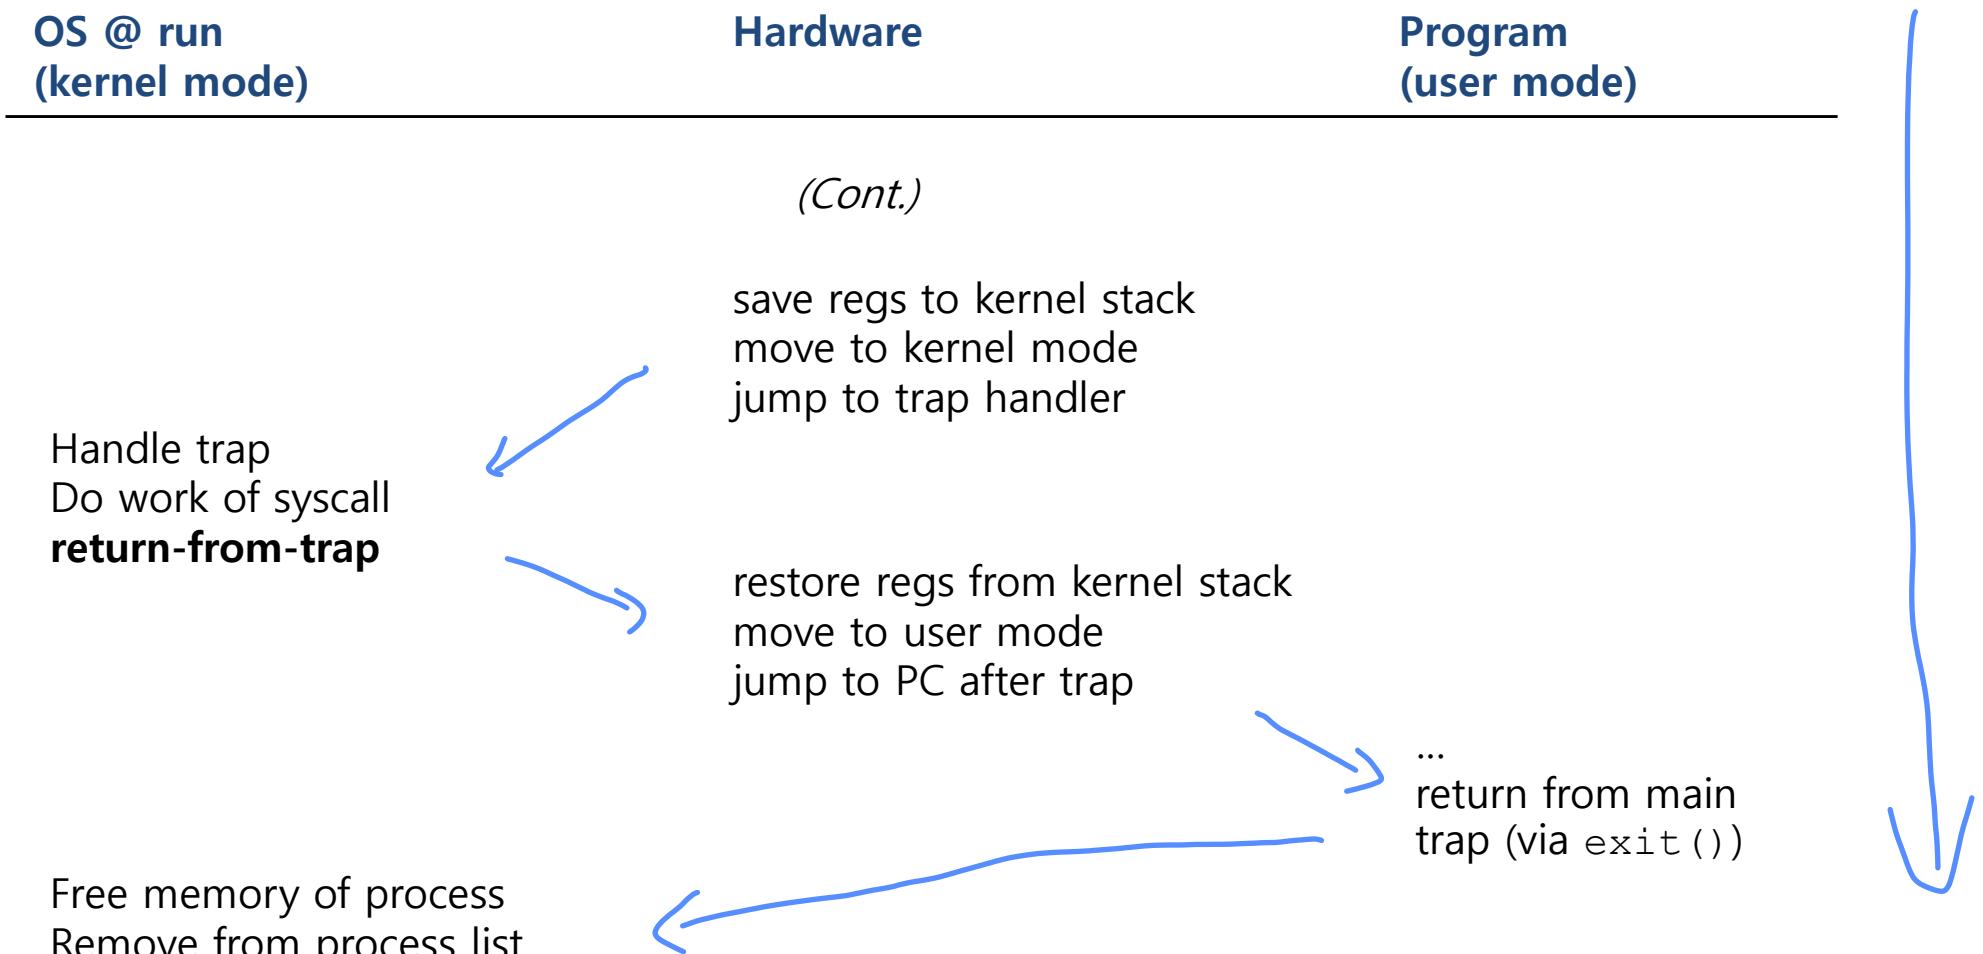
Program
(user mode)

restore regs from kernel stack
move to user mode
jump to main

Run main()
...
Call system
trap into OS

user stack :
kernel stack: system call
= 메모리 구조 .

Limited Direction Execution Protocol (Cont.)



Problem 2: Switching Between Processes

- How can the OS **regain control** of the CPU so that it can switch between *processes*?
 - A cooperative Approach: Wait for system calls
 - A Non-Cooperative Approach: The OS takes control

A cooperative Approach: Wait for system calls

프로세스가 자발적으로 yield라는 함수를 CPU를 내어줌. OS는 이를 감지함.

- Processes **periodically give up the CPU** by making system calls such as `yield`
 - The OS decides to run some other task
 - Application also transfer control to the OS **when they do something illegal**
 - Divide by zero
 - Try to access memory that it shouldn't be able to access
 - Ex) Early versions of the Macintosh OS,
The old Xerox Alto system

A process gets stuck in an infinite loop
→ **Reboot the machine**

yield로 안하고 illegal한 것들로 돌아온다.
한 번 더
가장 마지막.



A Non-Cooperative Approach: OS Takes Control

- A **timer interrupt**
 - During the boot sequence, the OS start the timer
 - The timer raise an interrupt every so many milliseconds
 - When the interrupt is raised :
 - The currently running process is halted
 - Save enough of the state of the program
 - A pre-configured interrupt handler in the OS runs

A **timer interrupt** gives OS the ability to run again on a CPU

Saving and Restoring Context

- **Scheduler** makes a decision:
 - Whether to continue running the current process, or switch to a different one
 - If the decision is made to switch, the OS executes context switch

Context Switch

- A low-level piece of assembly code
 - Save a few register values for the current process onto its kernel stack
 - General purpose registers
 - PC
 - kernel stack pointer
 - Restore a few for the soon-to-be-executing process from its kernel stack
 - Switch to the kernel stack for the soon-to-be-executing process

Limited Direction Execution Protocol (Timer interrupt)

OS @ boot
(kernel mode)

Hardware

initialize trap table

remember address of ...
syscall handler
timer handler

start interrupt timer

start timer
interrupt CPU in X ms

OS @ run
(kernel mode)

Hardware

Program
(user mode)

Process A

...

timer interrupt

save regs(A) to k-stack(A)
move to kernel mode
jump to trap handler

Limited Direction Execution Protocol (Timer interrupt)

OS @ run (kernel mode)	Hardware	Program (user mode)
<i>(Cont.)</i>		
Handle the trap		
Call switch() routine		
save regs(A) to proc-struct(A)		
restore regs(B) from proc-struct(B)		
switch to k-stack(B)		
return-from-trap (into B)		
	restore regs(B) from k-stack(B)	
	move to user mode	
	jump to B's PC	
		Process B
		...

Worried About Concurrency?

- What happens if, during interrupt or trap handling, another interrupt occurs?
- OS handles these situations:
 - Disable interrupts during interrupt processing 무시
 - Use a number of sophisticated locking schemes to protect concurrent access to internal data structures