


Operating System: CPU scheduling

Sang Ho Choi (shchoi@kw.ac.kr)

School of Computer & Information Engineering
KwangWoon University


CPU-I/O Burst Cycle

- Alternating sequence of CPU bursts and I/O bursts




CPU-I/O Burst Cycle

- Histogram of CPU burst times



CPU Scheduler

- CPU scheduler
 - selects a process from the processes that are ready to execute



- CPU scheduling decisions may take place when
 - (1) a process switches from running to waiting state (e.g. I/O request),
 - (2) a process switches from running to ready state (e.g. time slice expiration),
 - (3) a process switches from waiting to ready (e.g. I/O completion), or
 - (4) a process terminates

→ Scheduling under (1) and (4) is non-preemptive

→ Scheduling under (2) and (3) is preemptive

→ CPU 자원 자주 사용
반납.

→ 선점 → OS가 CPU 자원 배분 결정.



Scheduling Criteria



- CPU utilization

- keeps the CPU as **busy** as possible. (0% ~ 100%)



- Throughput

- The **number of processes** that are completed **per time unit**



- Turnaround time

- Time from the submission of a request **to time of completion**



- Waiting time

- Sum of time a process has been **waiting in the ready queue**



- Response time

- Time from the submission of a request **until the first response is produced**

수강료 청구!

Scheduling Criteria

- Scheduling algorithm's goals
 - Maximize the CPU utilization
 - Maximize the throughput
 - Minimize the turnaround time
 - Minimize the waiting time
 - Minimize the response time
 - It may be more important to minimize the variance than the average of response time
 - E.g. interactive system

하나가 놓아두면 다른거
남아질 수도 있어서
목표에 맞게 설계함.

평균보다 흔한이 적은거
중호(알고리즘 차원)

Scheduling: Introduction

- Workload assumptions:
 1. Each job runs for the same amount of time
 2. All jobs arrive at the same time
 3. All jobs only use the CPU (i.e., they perform no I/O)
 4. The run-time of each job is known

Scheduling Metrics


- Performance metric: **Turnaround time**
 - The time at which the job completes minus the time at which the job arrived in the system

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- Another metric is **fairness**
 - Performance and fairness are often at odds in scheduling

First In, First Out (FIFO)


- First Come, First Served (FCFS)
 - Very simple and easy to implement
- Example:
 - A arrived just before B which arrived just before C
 - Each job runs for 10 seconds



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why FIFO is not that great? - Convoy effect


- Let's relax assumption 1: Each job no longer runs for the same amount of time
- Example:
 - A arrived just before B which arrived just before C
 - A runs for 100 seconds, B and C run for 10 each



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$

Shortest Job First (SJF)


- Run the shortest job first, then the next shortest, and so on
 - Non-preemptive scheduler
- Example:
 - A arrived just before B which arrived just before C
 - A runs for 100 seconds, B and C run for 10 each



$$Average \ turnaround \ time = \frac{10 + 20 + 120}{3} = 50 \ sec$$

SJF with Late Arrivals from B and C

- Let's relax assumption 2: Jobs can arrive at any time
- Example:
 - A arrives at t=0 and needs to run for 100 seconds
 - B and C arrive at t=10 and each need to run for 10 seconds




$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$

Shortest Time-to-Completion First (STCF)

- Add **preemption** to SJF
 - Jobs are not available simultaneously (Assumption 2 relaxed)
 - Preemptive version of SJF (Assumption 3 relaxed)
 - Also known as Preemptive Shortest Job First (PSJF) or Shortest Remaining Time First (SRTF)
- A new job enters the system:
 - Determine of the remaining jobs and new job
 - Schedule the job which has the least time left

Shortest Time-to-Completion First (STCF)

- Example:
 - A arrives at t=0 and needs to run for 100 seconds
 - B and C arrive at t=10 and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

SJF Implementation

- Determining the lengths of next CPU bursts time is required
- Estimation by using the length of previous CPU burst times (**exponential averaging**)
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1 \rightarrow$ 적응의 속도
 4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$

SJF Implementation (Cont.)

- Next CPU burst time

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

- $\alpha = 0$

- $\tau_{n+1} = \tau_n$

- Recent history has no effect

- $\alpha = 1$

- $\tau_{n+1} = t_n$

- Only the most recent CPU burst matters


- If we expand the formula, we get

$$\begin{aligned}\tau_{n+1} &= \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots \\ &\quad + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ &\quad + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

- Since both α and $(1 - \alpha)$ are less than or equal to 1,
 - Each successive term has less weight than its predecessor

SJF Implementation (Cont.)

- Prediction of the length of the next CPU burst time



New scheduling metric: Response time

- The time from when the job arrives to the first time it is scheduled

$$T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$$

- STCF and related disciplines are not particularly good for response time

How can we build a scheduler that is
sensitive to response time?

Round Robin (RR) Scheduling

- Time slicing Scheduling
 - Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished
➤ Time slice is sometimes called a scheduling quantum
 - It repeatedly does so until the jobs are finished
 - The length of a time slice must be *a multiple of* the timer-interrupt period

RR is fair, but performs poorly on metrics such as turnaround time


response time은 3,21번

turnaround은 11번

인증자.


RR Scheduling Example

- A, B and C arrive at the same time
- They each wish to run for 5 seconds



SJF (Bad for Response Time)

$$T_{average \ response} = \frac{0 + 5 + 10}{3} = 5 \text{ sec}$$



RR with a time-slice of 1sec (Good for Response Time)

$$T_{average \ response} = \frac{0 + 1 + 2}{3} = 1 \text{ sec}$$

The length of the time slice is critical



- The shorter time slice
 - Better response time
 - The cost of context switching will dominate overall performance
- The longer time slice
 - Amortize the cost of switching
 - Worse response time

Deciding on the length of the time slice presents a **trade-off** to a system designer

Incorporating I/O

- Let's relax assumption 3: All programs perform I/O
- Example:
 - A and B need 50ms of CPU time each
 - A runs for 10ms and then issues an I/O request
 - I/Os each take 10ms
 - B simply uses the CPU for 50ms and performs no I/O
 - The scheduler runs A first, then B after

Incorporating I/O (Cont.)



Maximize the
CPU utilization

Overlap Allows Better Use of Resources

Incorporating I/O (Cont.)

- When a job initiates an I/O request
 - The job is blocked waiting for I/O completion
 - The scheduler should schedule another job on the CPU
- When the I/O completes
 - An interrupt is raised
 - The OS moves the process from blocked back to the ready state