

Operating System: Introduction to Paging

Sang Ho Choi (shchoi@kw.ac.kr)

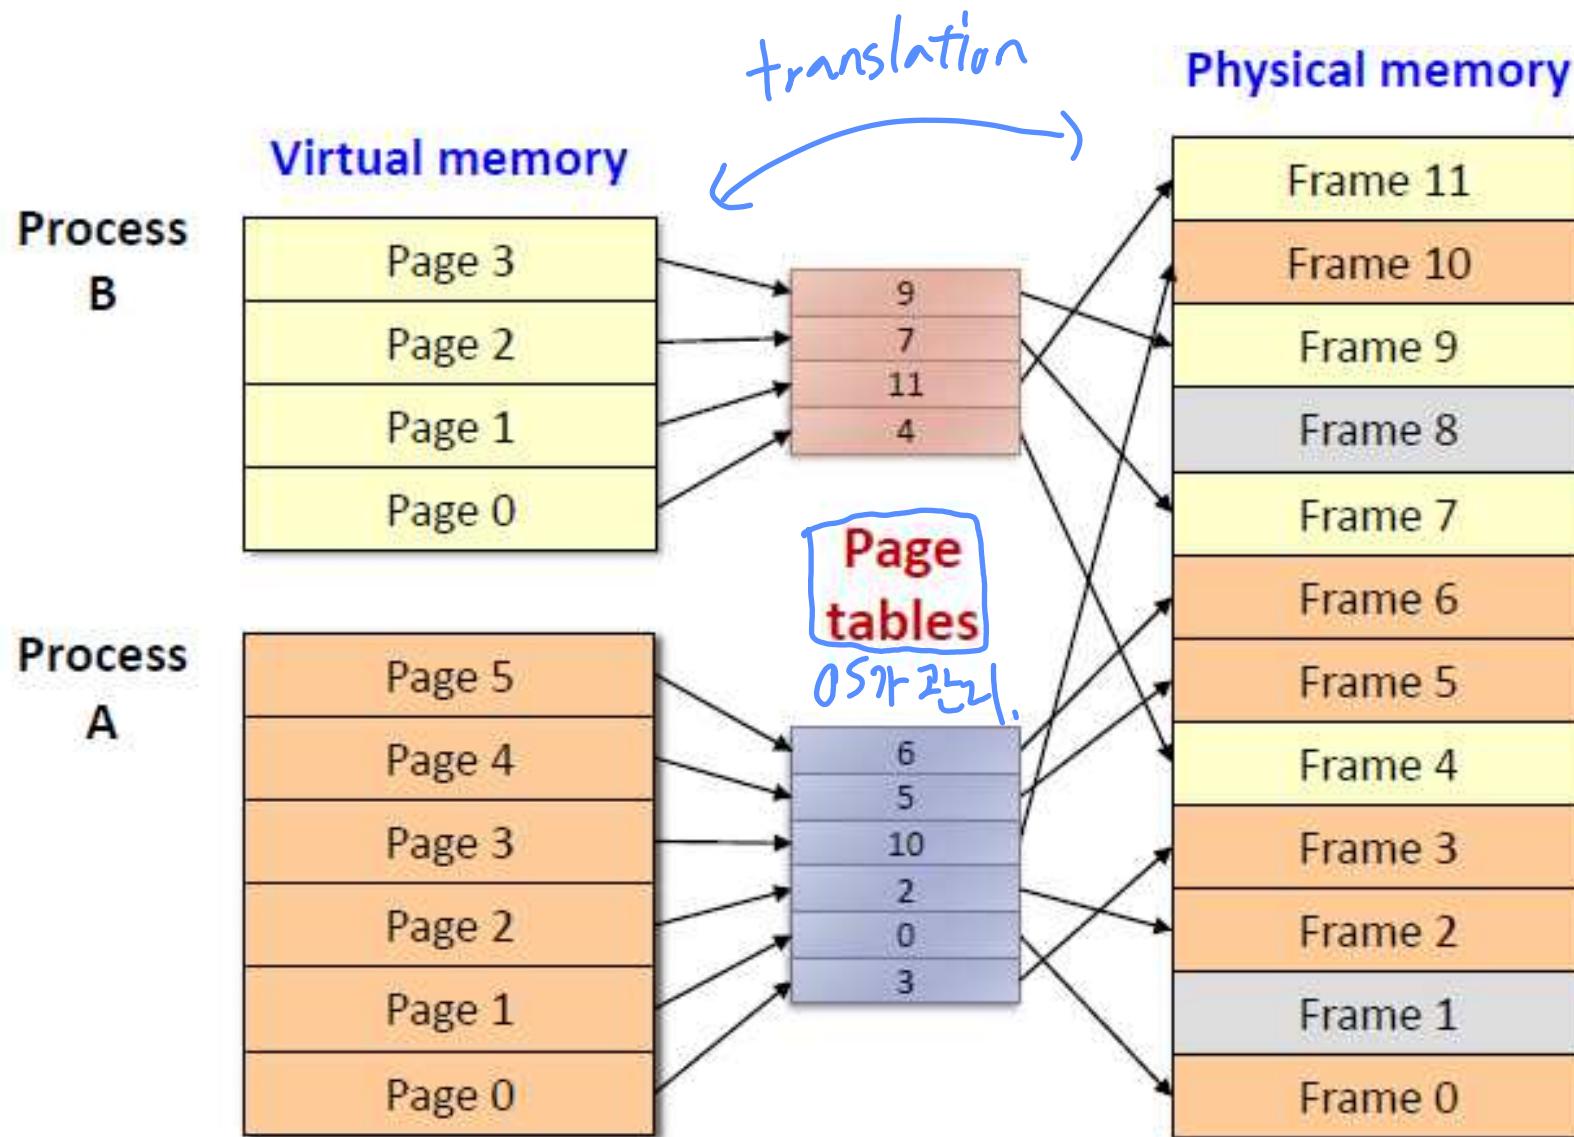
School of Computer & Information Engineering
KwangWoon University

Paging Overview

- Two approaches when solving most space-management problem
 - Segmentation: variable size of logical segments (code, stack, heap, etc.)
 - Fragmentation
 - Allocation becomes more challenging over time
 - **Paging** splits up address space into fixed-sized unit called a **page**
- With paging, physical memory is also split into some number of pages called a **page frame**
- **Page table** per process is needed to translate the virtual address to physical address

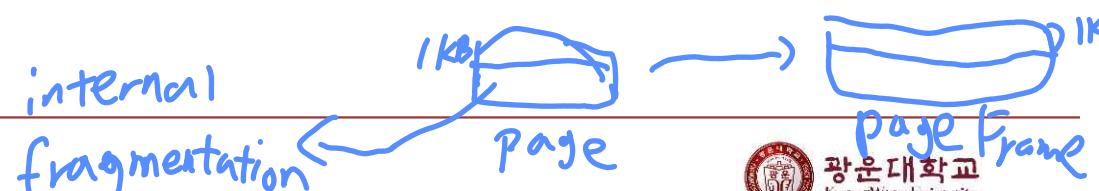
용어의(의)정

Paging Overview (Cont.)



Paging

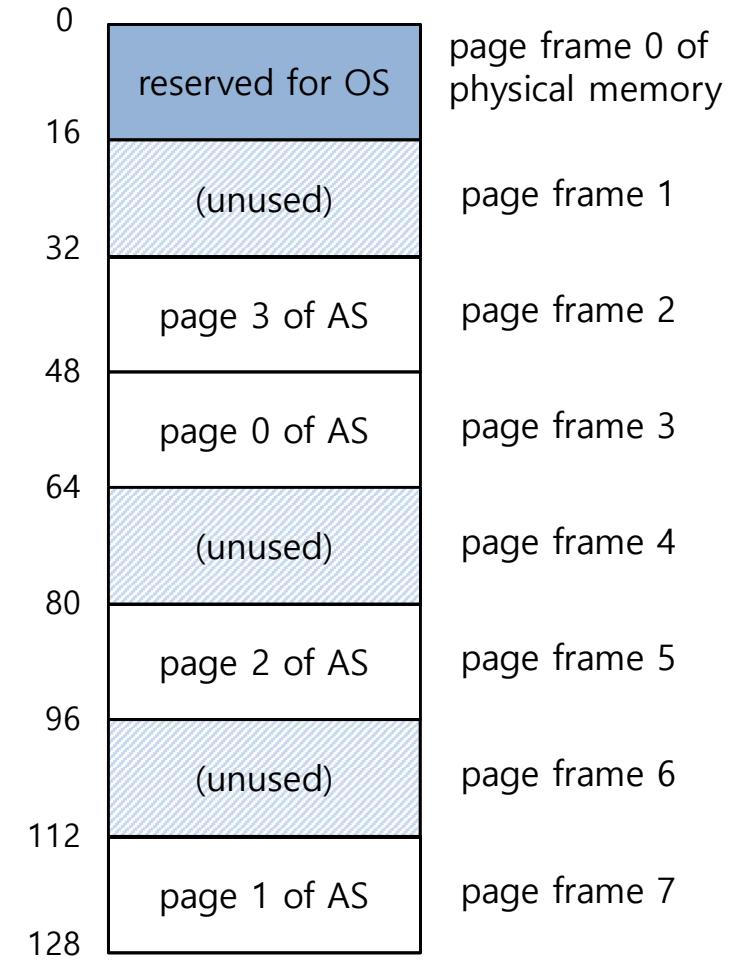
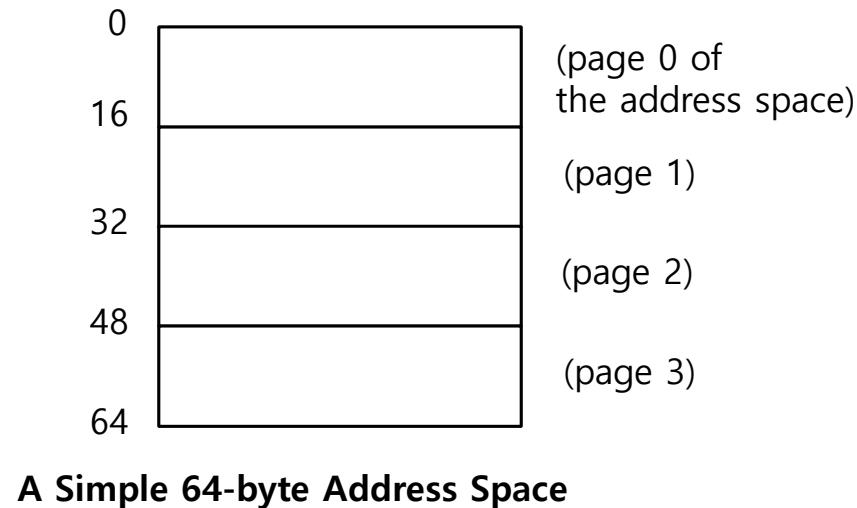
- Allows the physical address space of a process to be noncontiguous *프로세스 주소공간이 연속적이지 않음 X*
 - Divide virtual memory into blocks of same size (**pages**)
 - Divide physical memory into fixed-size blocks (**frames**)
 - Page (or frame) size is power of 2 (typically 512B - 8KB) *대부분 이정도 사이즈.*
 - Eases memory management
 - OS keeps track of all free frames
 - To run a program of size n pages, need to find n free frames and load the program
 - Set up a page table to translate virtual to physical addresses
 - No external fragmentation

*기존 방식은 프로그램공간을 위한
알맞는 공간을 hole에서 찾아야 했으나.
그걸 필요 X*
- 
- internal fragmentation*
- Page Frame*
- 1KB*
- 1KB*



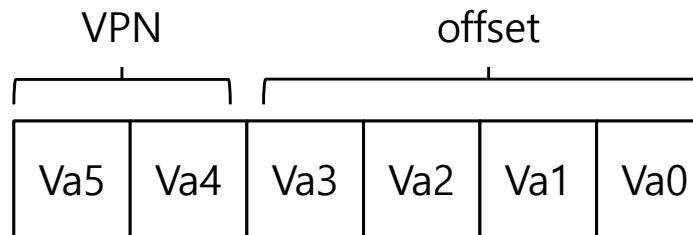
A Simple Paging Example

- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages

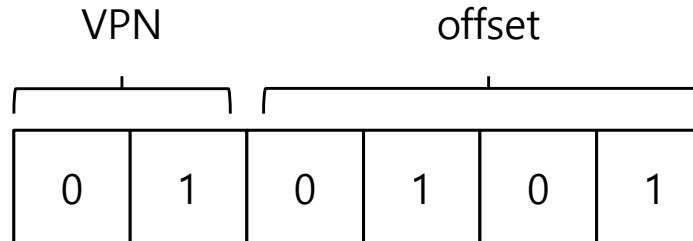


Address Translation

- Two components in the virtual address
 - VPN: virtual page number
 - Offset: offset within the page

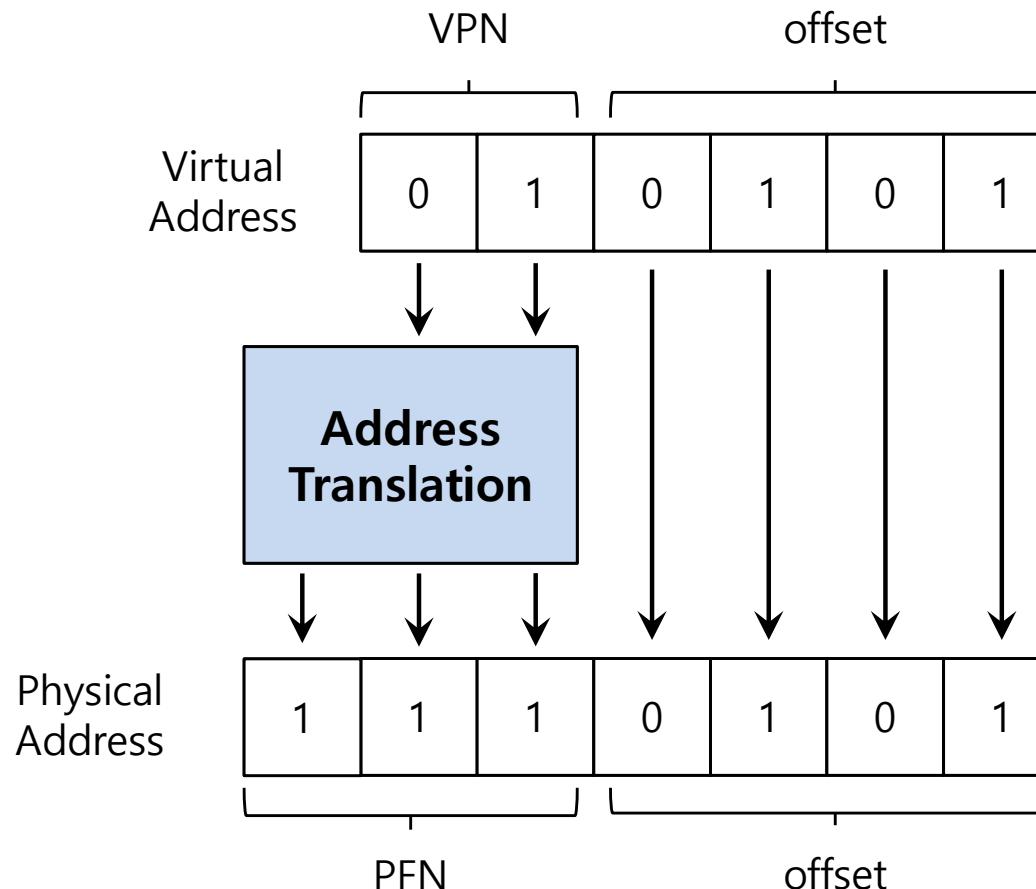


- Example: virtual address 21 in 64-byte address space



Address Translation (Cont.)

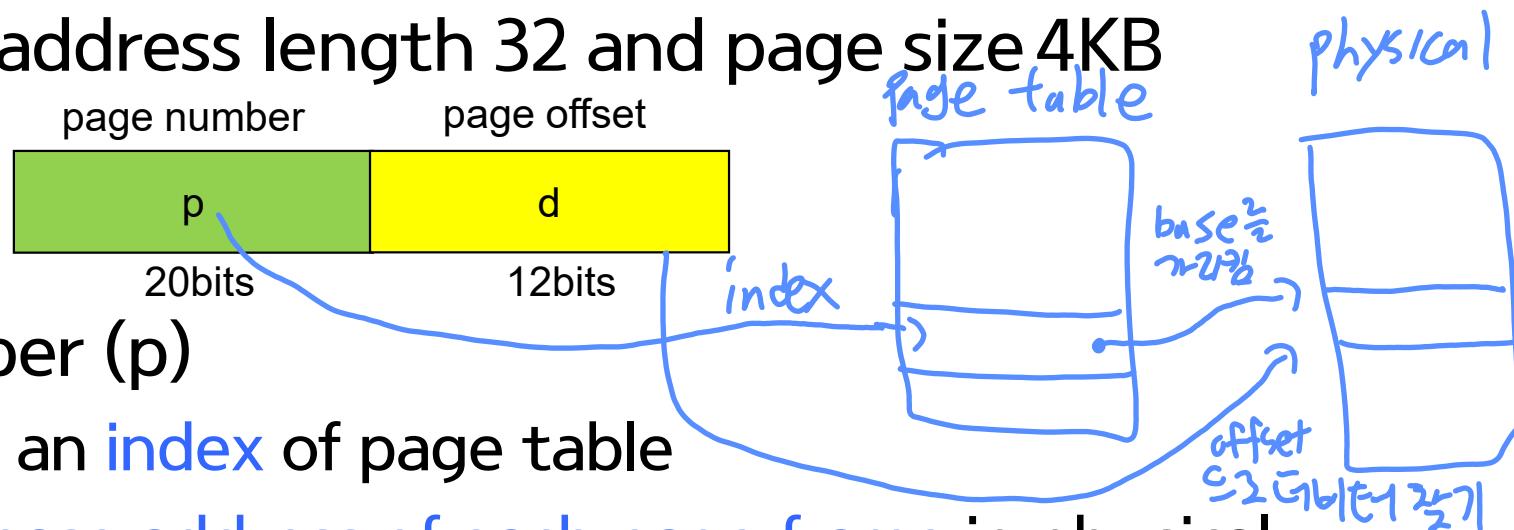
- The virtual address 21 in 64-byte address space



page
Frame Number

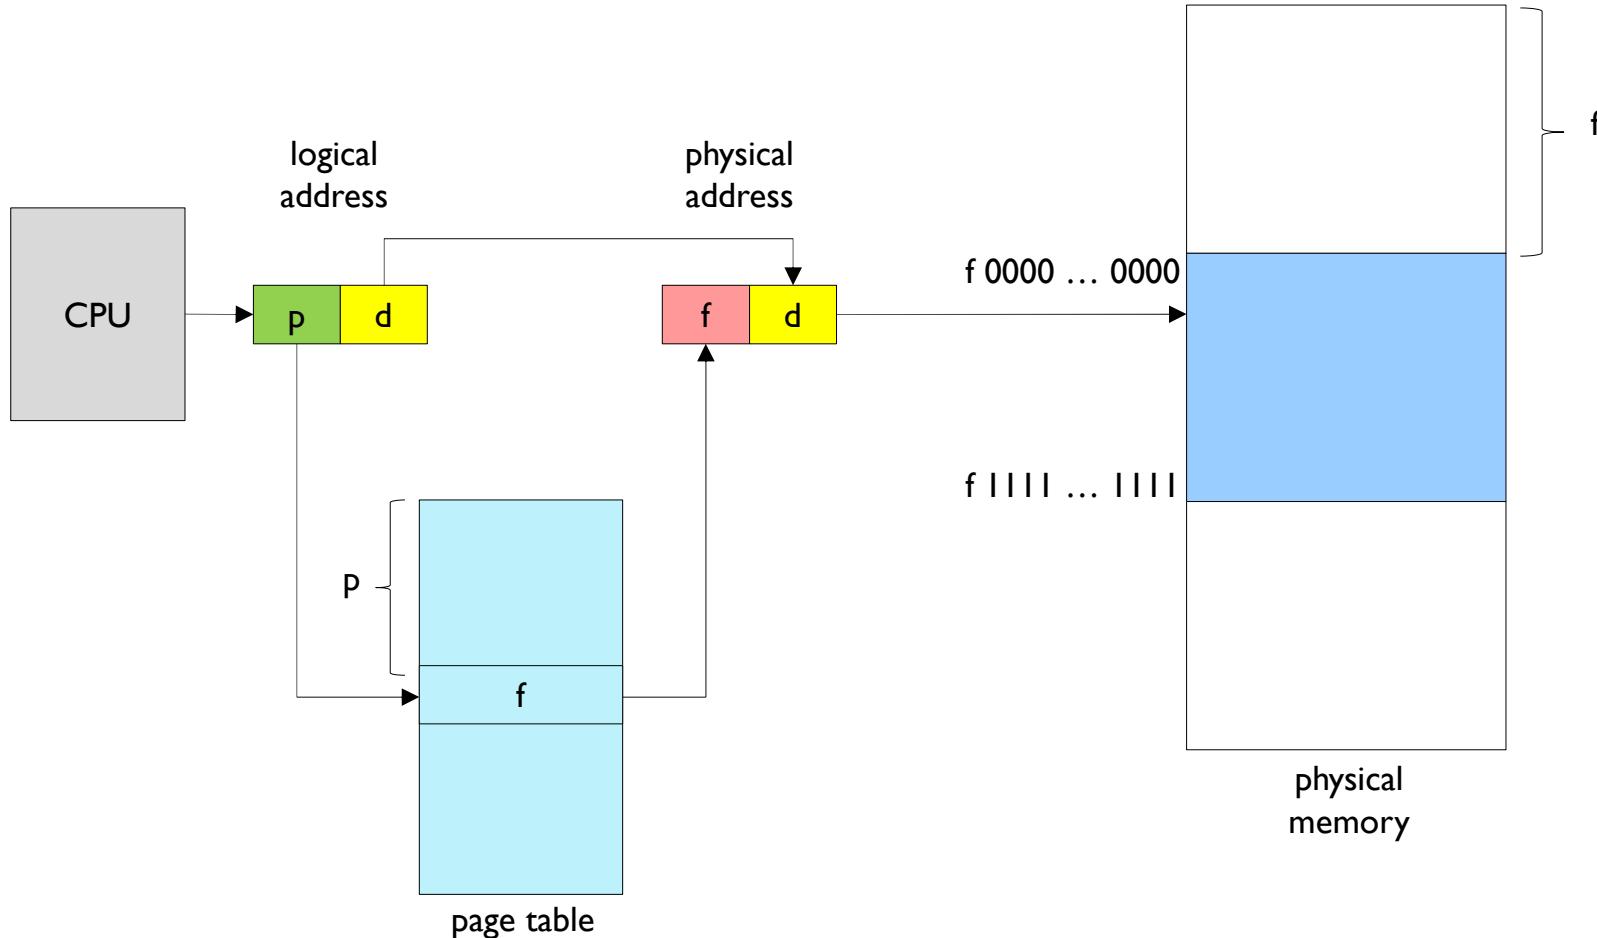
Address Translation (Cont.)

- Another example
 - The logical address generated by CPU is divided into two parts
 - E.g. Given address length 32 and page size 4KB
 - Page number (p)
 - Is used as an **index** of page table
 - Contains **base address** of each page frame in physical memory
 - Page offset (d)
 - Is combined with base address to define the physical memory address



Address Translation (Cont.)

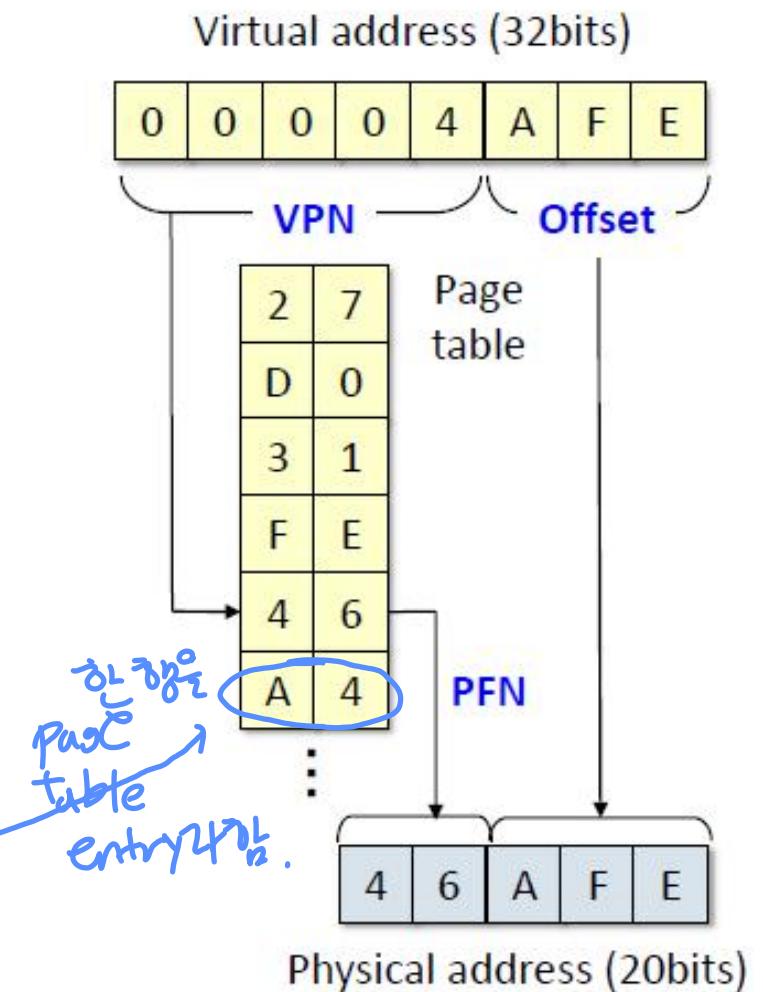
- Address Translation Scheme



Address Translation (Cont.)

- Example
 - Virtual address: 32 bits
 - Physical address: 20 bits
 - Page size: 4KB
 - Offset: 12 bits
 - VPN: 20 bits
 - Page table entries: 2^{20}
 - Page table size?

= 1 MB

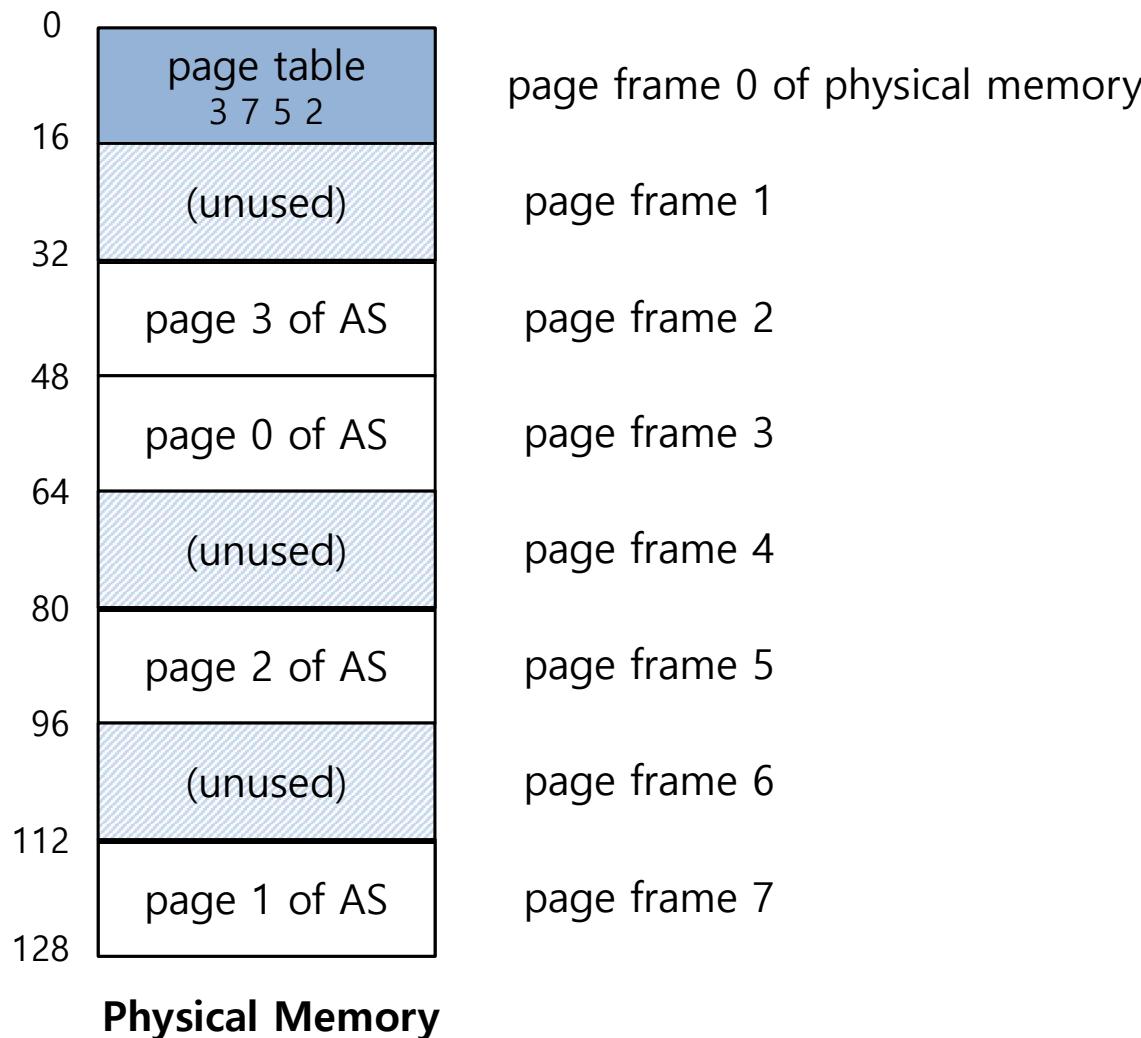


Paging: Spatial Overhead

- Where Are Page Tables Stored?
 - Page tables for each process are stored in memory
각 프로세스의 푸터는 메모리에 저장됩니다. 단위는 각 프로세스입니다.
- Page tables can get awfully large
 - 32-bit address space with 4-KB pages, 20 bits for VPN
 $\triangleright 4M B = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$

Page Table in Kernel Physical Memory

- Page tables for each process are stored in memory



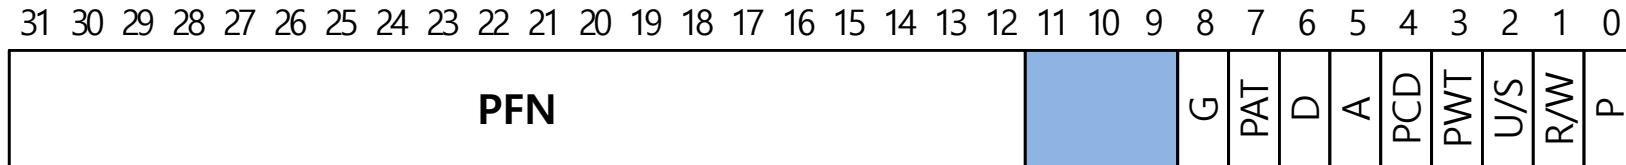
What Is In The Page Table?

- The page table is just a data structure that is used to map the virtual address to physical address
 - Simplest form: a linear page table, an array
- The OS indexes the array by VPN, and looks up the page-table entry

Common Flags of Page Table Entry

- Valid Bit
 - Indicates whether the particular translation is valid
 - It is checked each time a virtual address is used
- Protection Bit
 - Indicates whether the page could be read from, written to, or executed from
- Present Bit
 - Indicates whether this page is in physical memory or on disk(swapped out)
- Dirty Bit
 - Indicates whether the page has been modified since it was brought into memory
- Reference Bit (Accessed Bit)
 - Indicates that a page has been accessed

Example: x86 Page Table Entry



An x86 Page Table Entry(PTE)

- P: present
- R/W: read/write bit
- U/S: supervisor
- A: accessed bit
- D: dirty bit
- PFN: the page frame number

Paging: Temporal Overhead

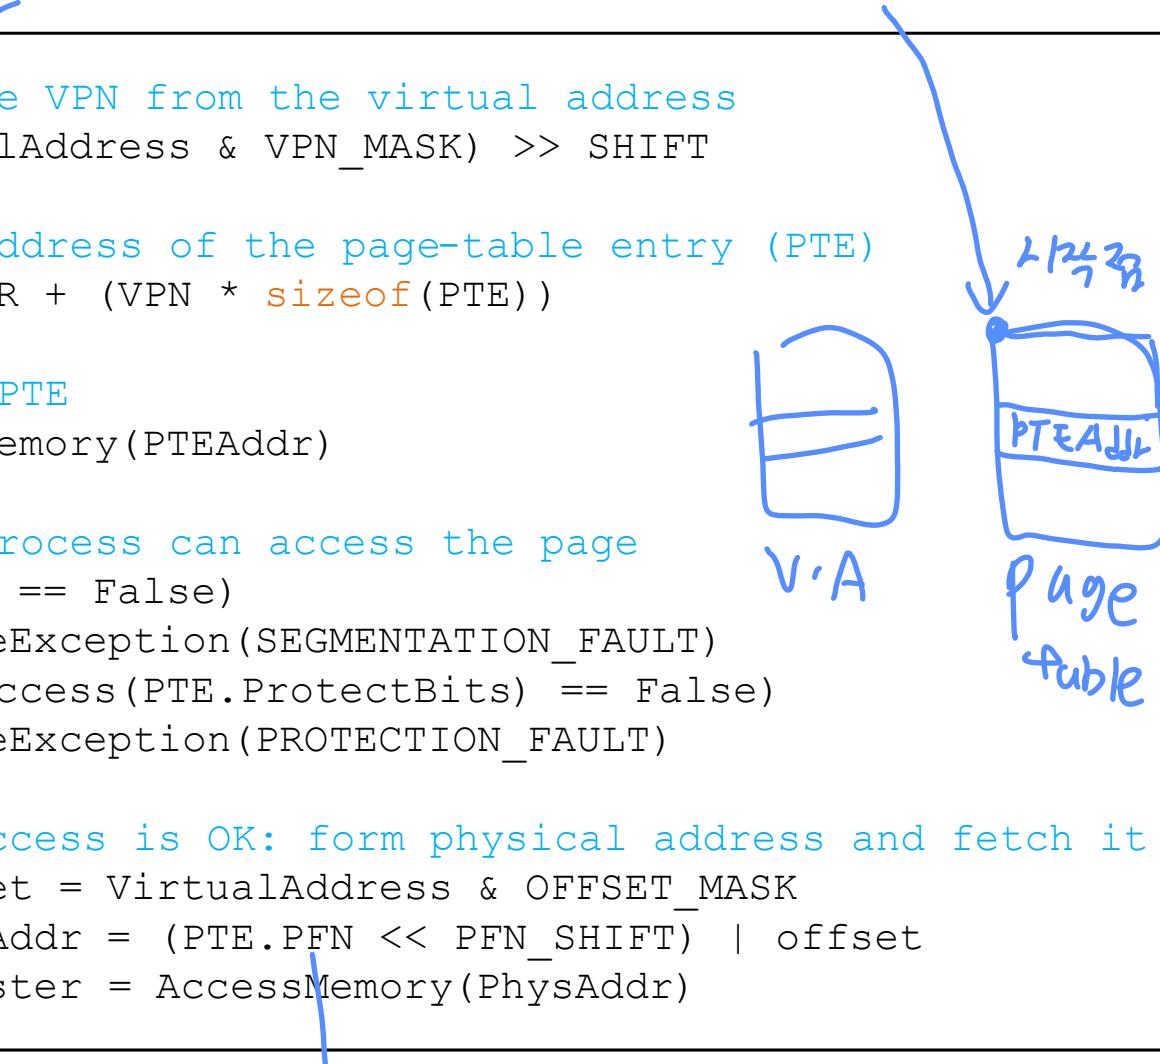
- To find a location of the desired PTE, the starting location of the page table is needed
- For every memory reference, paging requires the OS to perform one extra memory reference

Page table 또한 주소를 찾기 위해
이를 참조하는데 추가적인 주소를 찾는다.

Accessing Memory With Paging

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PTE)
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```

page table base register



page frame number.

Demand Paging

- Divide memory and image into fixed size, respectively
 - page frame and page
 - Size is power of 2, between 512B and 8KB
- Entire program image resides **on disk**
- When the program starts, just **load the first page only**
- The **rest** of pages are loaded in memory **on-demand**
- A particular page X of the program can be either
 - already loaded **in memory** page frame Y, or
 - never been loaded before, it is **in disk**
- Pages can be placed **anywhere** in memory
- Whenever CPU presents an address, MMU looks up **page table**
 - For translating logical address to physical address

disk에 일단 넣고
그 뒤 페이지만
필수화시 쓰.

Memory management Unit

Paging: Pros

- No external fragmentation
- Fast to allocate and free
 - A list or bitmap for free page frames
 - Allocation: no need to find contiguous free space
 - Free: no need to coalesce with adjacent free space
- Easy to “page out” portions of memory to disk
파일로 잡아내거나 디스크로 보내는거
 - Page size is chosen to be a multiple of disk block sizes
 - Use valid bit to detect reference to “paged-out” pages
 - Can run process when some pages are on disk
- Easy to protect and share pages

Paging: Pros (Cont.)

클래스

- Shared pages

code 1
code 2
code 3
data 1

process P₁

3
4
6
1

Page table
for P₁

code 1
code 2
code 3
data 2

process P₂

3
4
6
7

Page table
for P₂

code 1
code 2
code 3
data 3

process P₃

3
4
6
2

Page table
for P₃

0
1
2
3
4
5
6
7
8
9
10
11



Paging: Cons

- Internal fragmentation
 - Wasted memory grows with larger pages
- Memory reference overhead (temporal overhead)
 - Doubles the number memory references per instruction
 - Solution: get hardware support (**TLBs**)
- Storage needed for page tables (spatial overhead)
 - Needs one PTE for each page in virtual address space
 - 32-bit address space with 4KB page size: 2^{20} PTEs
 - 4 bytes/PTE: 4MB per page table
 - 100 processes in the system: total 400MB of page tables
 - Solution: store valid PTEs only