

# Operating System: Semaphores

---

Sang Ho Choi ([shchoi@kw.ac.kr](mailto:shchoi@kw.ac.kr))

School of Computer & Information Engineering  
KwangWoon University

# Semaphore: A definition

- Semaphore
  - An object **with an integer value, S**
- We can manipulate with two routines:
  - sem\_wait(S) *lock 가능한*
  - originally called P(S)
  - sem\_post(S)  
➤ originally called V(S) *unlock //*

```
1  sem_t m;  
2  sem_init(&m, 0, X);  
3  
4  sem_wait(&m);  
5  //critical section here  
6  sem_post(&m);
```

Using semaphore

# Semaphore: Initialization

- Initialization

```
1 #include <semaphore.h>
2 sem_t s;
3 sem_init(&s, 0, 1); // initialize s to the value 1
```

한 프로세스 내

스레드들 사이에 서로 다른 경우

- Declare a semaphore `s` and initialize it to the value 1
- The second argument, 0, indicates that the semaphore is shared between *threads in the same process*

# Semaphore: Interact with semaphore

- `sem_wait()`

```
1 int sem_wait(sem_t *s) {  
2     decrement the value of semaphore s by one  
3     wait if value of semaphore s is negative  
4 }
```

값이 감소  
시작으로 이동할 수 있는지 확인.

- If the value of the semaphore was *one or higher* when called `sem_wait()`, return right away 마지막
- It will cause the caller to suspend execution waiting for a subsequent post 포스트 흐름 차단
- When negative, the value of the semaphore is equal to the number of waiting threads

즉, 그 숫자 만큼 기다리고 있는 thread의 수

# Semaphore: Interact with semaphore (Cont.)

- `sem_post()`

```
1 int sem_post(sem_t *s) {  
2     increment the value of semaphore s by one  
3     if there are one or more threads waiting, wake one  
4 }
```

- Simply increments the value of the semaphore
- If there is a thread waiting to be woken, wakes one of them up

값 증가

터미네이터는 큐에서 하나를 꺼내기

# Binary Semaphores (Locks)

- What should  $x$  be?
  - The initial value should be 1

```
1 sem_t m;  
2 sem_init(&m, 0, X); // initialize semaphore to X; what should X be?  
3  
4 sem_wait(&m);  
5 //critical section here  
6 sem_post(&m);
```

binary semaphore은 1개만 같은 순간에 하나가  
 $X$ 가 1개여야 한다고?

# Thread Trace: Single Thread Using a Semaphore

Value of Semaphore	Thread 0	Thread 1
1		
1	call sem_wait()	
0	sem_wait() returns	
0	(crit sect)	
0	call sem_post()	
1	sem_post() returns	

설명하기

# Thread Trace: Two Threads Using a Semaphore

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() retruns	Running		Ready
0	(crit set: begin)	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem < 0) → sleep	sleeping
-1		Running	Switch → T0	sleeping
-1	(crit sect: end)	Running		sleeping
-1	call sem_post()	Running		sleeping
0	increment sem	Running		sleeping
0	wake(T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	sem_wait() retruns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

# Semaphores As Condition Variables

```
1 sem_t s;
2
3 void *
4 child(void *arg) {
5     printf("child\n");
6     sem_post(&s); // signal here: child is done
7     return NULL;
8 }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

A Parent Waiting For Its Child

– What should x be?

➤ The value of semaphore should be set to is 0

아래와 같이  
작동하는 것을  
알수 있다.  
Condition Variable  
는 대체로 다음과 같은  
방법으로 사용된다.

parent: begin  
child  
parent: end

The execution result

# Thread Trace: Parent Waiting For Child (Case 1)

- The parent call `sem_wait()` before the child has called `sem_post()`

Value	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	call <code>sem_wait()</code>	Running		Ready
-1	decrement sem	Running		Ready
-1	$(\text{sem} < 0) \rightarrow \text{sleep}$	sleeping		Ready
-1	<i>Switch→Child</i>	sleeping	child runs	Running
-1		sleeping	call <code>sem_post()</code>	Running
0		sleeping	increment sem	Running
0		Ready	wake(Parent)	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	<i>Interrupt; Switch→Parent</i>	Ready
0	<code>sem_wait()</code> retruns	Running		Ready

# Thread Trace: Parent Waiting For Child (Case 2)

- The child runs to completion before the parent  
`call sem_wait()`

Value	Parent	State	Child	State
0	Create (Child)	Running	(Child exists; is runnable)	Ready
0	<i>Interrupt; switch→Child</i>	Ready	child runs	Running
0		Ready	call sem_post()	Running
1		Ready	increment sem	Running
1		Ready	wake (nobody)	Running
1		Ready	sem_post() returns	Running
1	parent runs	Running	<i>Interrupt; Switch→Parent</i>	Ready
1	call sem_wait()	Running		Ready
0	decrement sem	Running		Ready
0	(sem<0)→awake	Running		Ready
0	sem_wait() retruns	Running		Ready

Case 1, 2 모두 child가 먼저 실행되는 경우 같은 결과.

interrupt가 일어나거나 다른 경우 Case 1과 같은 결과를 얻어야 한다.

# Classical Problems of Synchronization

---

- Solutions with semaphore for the following problems
  - Producer/Consumer (Bounded buffer) Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem

Semaphore  
Semaphore  
Semaphore

# Bounded buffer problem

- Producer
    - Produce data items
    - Wish to place data items in a buffer
  - Consumer
    - Grab data items out of the buffer consume them in some way
  - Example: Multi-threaded web server
    - A producer puts HTTP requests in to a work queue
    - Consumer threads take requests out of this queue and process them
- buffer의 데이터를 사용할 때  
제한
- 내가 데이터를 가져갈 수 있는  
방법
- 나에게 전송하기  
할 수 있는  
방법.

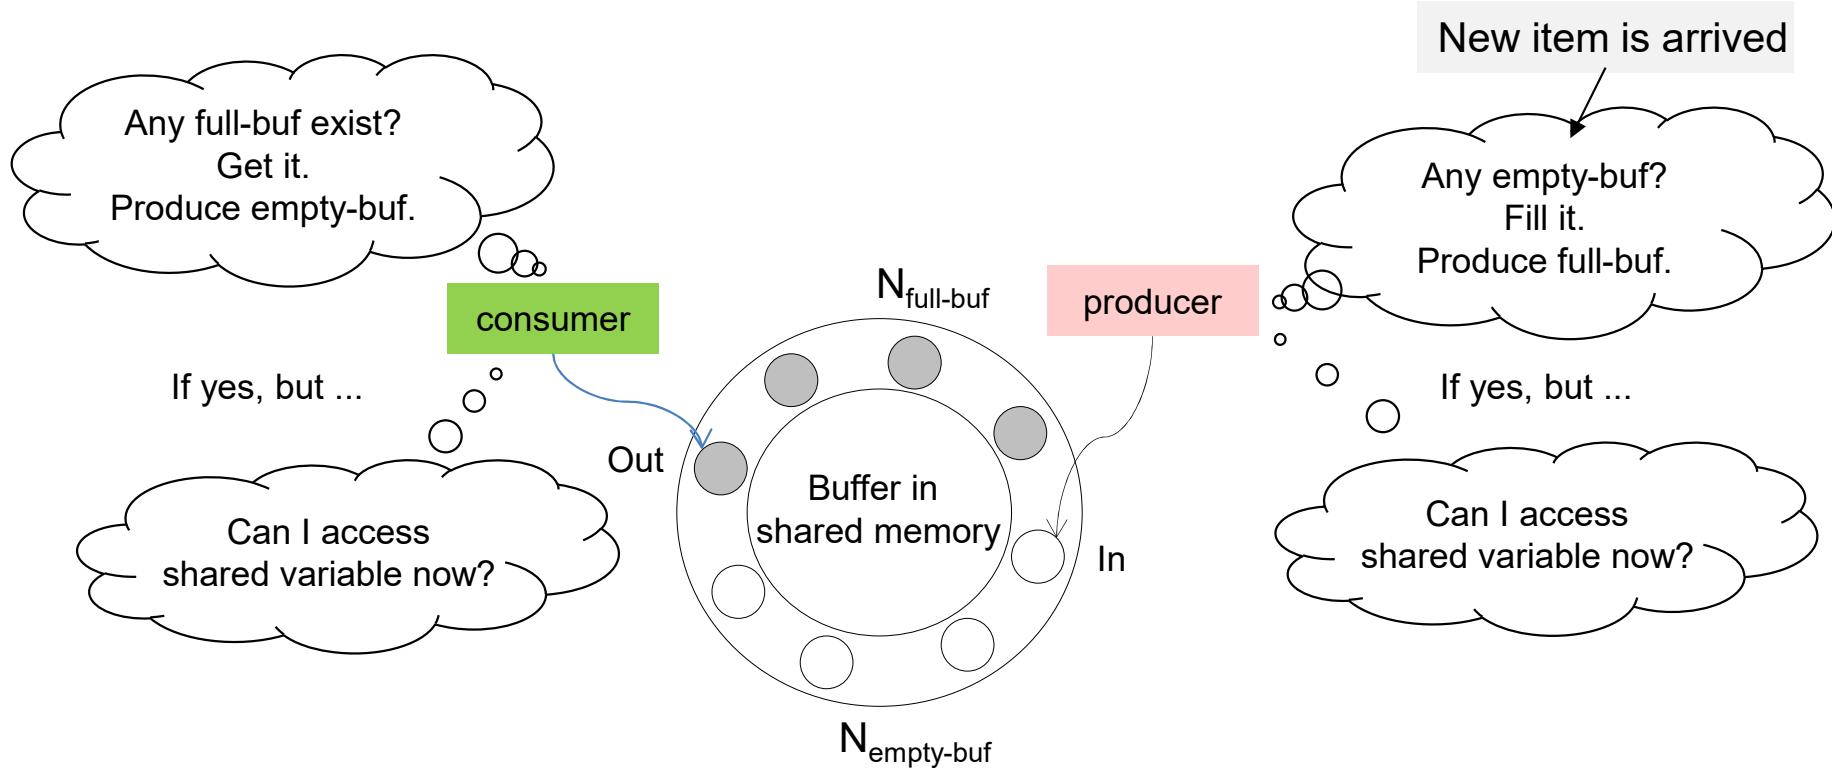
# Bounded buffer problem (Cont.)

- A bounded buffer is used when you pipe the output of one program into another
  - Example: `grep foo file.txt | wc -l`
    - The grep process is the producer
    - The wc process is the consumer
    - Between them is an in-kernel bounded buffer
  - Bounded buffer is Shared resource → Synchronized access is required

file.txt에서 foo라는 애  
찾아서 그 출력을 wc  
로 넘기는데 그 사이  
파이프로 넘길때  
bounded buffer라는  
걸 사용해 그래서 이  
를 bounded buffer는  
공유자원 grep은  
producer wc는  
consumer로 비유함

# Bounded buffer problem (Cont.)

- Supposed that there are N buffers



# A Solution

- Producer: put () interface
  - Wait for a buffer to become *empty* in order to put data into it
- Consumer: get () interface
  - Wait for a buffer to become *filled* before using it

```
1 int buffer[MAX];
2 int fill = 0;
3 int use = 0;
4
5 void put(int value) {
6     buffer[fill] = value;      // line f1
7     fill = (fill + 1) % MAX;  // line f2
8 }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;   // line g2
13     return tmp;
14 }
```

# A Solution (Cont.)

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);           // line P1
8          put(i);                  // line P2
9          sem_post(&full);         // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);        // line C1
17         tmp = get();            // line C2
18         sem_post(&empty);       // line C3
19         printf("%d\n", tmp);
20     }
21 }
22 ...
```

공유자원의 접근을 서로 제어해주니까 producer가 full을 늘려주기 전엔 consumer가 접근을 못하고 consumer가 empty를 늘려주기 전까진 producer가 접근을 못 함 따라서 mutual exclusive가 잘된다.

## First Attempt: Adding the Full and Empty Conditions

# A Solution (Cont.)

```
21 int main(int argc, char *argv[]) {  
22     // ...  
23     sem_init(&empty, 0, MAX);           // MAX buffers are empty to begin with...  
24     sem_init(&full, 0, 0);            // ... and 0 are full  
25     // ...  
26 }
```

## First Attempt: Adding the Full and Empty Conditions (Cont.)

- Imagine that MAX is greater than 1
  - If there are multiple producers, race condition can happen at line f1
  - It means that the old data there is overwritten
- We've forgotten here is mutual exclusion
  - The filling of a buffer and incrementing of the index into the buffer is a critical section

하지만 producer 혹은 consumer가 여러개일때는 MAX값이 2이상이라 상호배제가 잘 되지 않을 것이다.

# A Solution: Adding Mutual Exclusion

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&mutex);           // line p0 (NEW LINE)
9          sem_wait(&empty);         // line p1
10         put(i);                 // line p2
11         sem_post(&full);        // line p3
12         sem_post(&mutex);       // line p4 (NEW LINE)
13     }
14 }
15
(Cont.)
```

## Adding Mutual Exclusion (Incorrectly)

# A Solution: Adding Mutual Exclusion

(Cont.)

```
16 void *consumer(void *arg) {  
17     int i;  
18     for (i = 0; i < loops; i++) {  
19         sem_wait(&mutex);           // line c0 (NEW LINE)  
20         sem_wait(&full);          // line c1  
21         int tmp = get();          // line c2  
22         sem_post(&empty);          // line c3  
23         sem_post(&mutex);          // line c4 (NEW LINE)  
24         printf("%d\n", tmp);  
25     }  
26 }
```

## Adding Mutual Exclusion (Incorrectly)

뮤텍스를 추가했으나  
이렇게 되면 한쪽은  
뮤텍스가 해제되길 기  
다리는데 다른 한쪽은  
세마포어 값을 증가시  
키길 기다리게됨  
deadlock이 생기는거  
지

# A Solution: Adding Mutual Exclusion (Cont.)

- Imagine two threads: one producer and one consumer
  - The consumer acquires the `mutex` (line c0)
  - The consumer calls `sem_wait()` on the full semaphore (line c1)
  - The consumer is blocked and yields the CPU
    - The consumer still holds the mutex!
  - The producer calls `sem_wait()` on the binary `mutex` semaphore (line p0)
  - The producer is now **stuck** waiting too. **a classic deadlock**

```
8 sem_wait(&mutex);           // line p0
9 sem_wait(&empty);          // line p1
10 put(i);                  // line p2
11 sem_post(&full);          // line p3
12 sem_post(&mutex);          // line p4
```

producer

```
19 sem_wait(&mutex);          // line c0
20 sem_wait(&full);           // line c1
21 int tmp = get();           // line c2
22 sem_post(&empty);          // line c3
23 sem_post(&mutex);          // line c4
```

consumer

# Finally, A Working Solution

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&empty);           // line p1
9          sem_wait(&mutex);         // line p1.5 (MOVED MUTEX HERE...)
10         put(i);                // line p2
11         sem_post(&mutex);        // line p2.5 (... AND HERE)
12         sem_post(&full);         // line p3
13     }
14 }
15
(Cont.)
```

## Adding Mutual Exclusion (Correctly)

# Finally, A Working Solution

그래서 뮤텍스를 버퍼에 값을 넣기 바로 직전에 삽입하였음

```
(Cont.)  
16 void *consumer(void *arg) {  
17     int i;  
18     for (i = 0; i < loops; i++) {  
19         sem_wait(&full);           // line c1  
20         sem_wait(&mutex);        // line c1.5 (MOVED MUTEX HERE...)  
21         int tmp = get();         // line c2  
22         sem_post(&mutex);        // line c2.5 (... AND HERE)  
23         sem_post(&empty);        // line c3  
24         printf("%d\n", tmp);  
25     }  
26 }  
27  
28 int main(int argc, char *argv[]) {  
29     // ...  
30     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with ...  
31     sem_init(&full, 0, 0);   // ... and 0 are full  
32     sem_init(&mutex, 0, 1); // mutex=1 because it is a lock  
33     // ...  
34 }
```

## Adding Mutual Exclusion (Correctly)

# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – **only read** the data set
  - Writers – **read and write** the data set
- Problem
  - **Multiple readers can read at the same time**
  - **Only one single writer can access** the shared data at the same time
- Shared Data
  - data set
  - integer **readcount**: the number of readers
  - semaphore **mutex**: mutual exclusion for access to read count
  - semaphore **wrt**: mutual exclusion for writer

reader는 여러명이 공유자원에 접근해도 상관없지만 writer가 문제 그래서 이를 해결해보자

# Readers-Writers Problem (Cont.)

- The structures of writer and reader
  - semaphore mutex = 1, wrt = 1;
  - int readcount = 0;

누군가 읽고 있을때 write를 못하게끔 wrt를 감소시키고 또 누군가가 쓰고 있을때에도 write를 못하게끔 wrt를 감소시킨다. 또 다수가 reading을 할때 readcount 값을 증가시키는데 있어서 경쟁조건을 해소시키기 위해 mutex를 추가한 것

writer

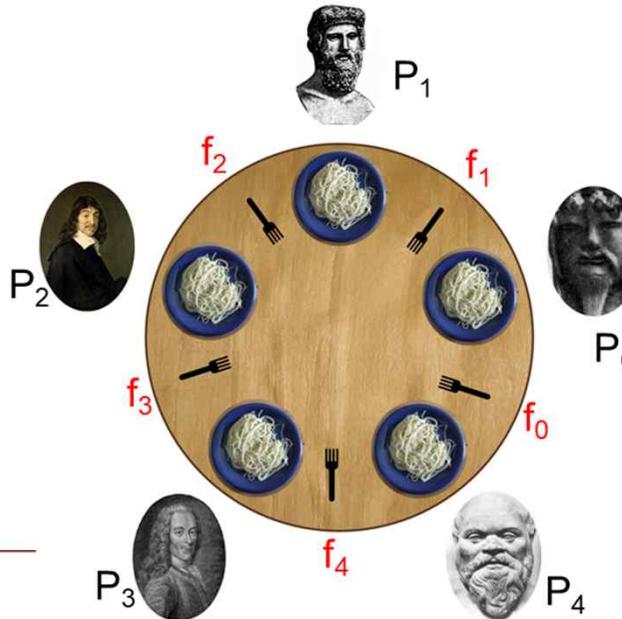
```
while (true) {  
    sem_wait(wrt);  
  
    // writing is performed  
  
    sem_post(wrt);  
}
```

reader

```
while (true) {  
    sem_wait(mutex);  
    readcount++;  
    if (readcount == 1)    sem_wait(wrt);  
    sem_post(mutex);  
  
    // reading is performed  
  
    sem_wait (mutex);  
    readcount--;  
    if (readcount == 0)    sem_post(wrt);  
    sem_post(mutex);  
}
```

# Dining Philosophers Problem

- Assume there are five “philosophers” sitting around a table
  - Between each pair of philosophers is a single fork (five total)
  - The philosophers each have times where they think, and don’t need any forks, and times where they eat
  - In order to *eat*, a philosopher needs **two forks**, both the one on their *left* and the one on their *right*
  - The contention for these forks



# Dining Philosophers (Cont.)

- Key challenge
  - There is no deadlock
  - No philosopher starves and never gets to eat
  - Concurrency is high

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

Basic loop of each philosopher

```
// helper functions  
int left(int p) { return p; }  
  
int right(int p) {  
    return (p + 1) % 5;  
}
```

Helper functions (Downey's solutions)

- Philosopher  $p$  wishes to refer to the fork on their left → call  $\text{left}(p)$
- Philosopher  $p$  wishes to refer to the fork on their right → call  $\text{right}(p)$

# Dining Philosophers (Cont.)

- We need some semaphore, one for each fork:

```
sem_t forks[5]
```

```
1 void getforks() {
2     sem_wait(forks[left(p)]);
3     sem_wait(forks[right(p)]);
4 }
5
6 void putforks() {
7     sem_post(forks[left(p)]);
8     sem_post(forks[right(p)]);
9 }
```

이런 경우에는 deadlock이 발생 가능 5명이 모두 왼쪽 포크를 얻은 상황에서 모두가 오른쪽 포크를 얻으려고 한다면 누구도 얻을 수 없을 것

The `getforks()` and `putforks()` Routines (Broken Solution)

- Deadlock occur!

- If each philosopher happens to grab the fork on their left before any philosopher can grab the fork on their right
- Each will be stuck *holding one fork* and waiting for another, *forever*

# A Solution: Breaking The Dependency

- Change how forks are acquired
  - Let's assume that philosopher 4 acquire the forks in a *different order*

```
1 void getforks() {  
2     if (p == 4) {  
3         sem_wait(forks[right(p)]);  
4         sem_wait(forks[left(p)]);  
5     } else {  
6         sem_wait(forks[left(p)]);  
7         sem_wait(forks[right(p)]);  
8     }  
9 }
```

간단하게 4명이 왼쪽 포크를 얻었다면 마지막 한명은 오른쪽 포크부터 얻도록 순서를 약간 조정하면 deadlock문제를 해결 가능

- There is no situation where each philosopher grabs one fork and is stuck waiting for another. The cycle of waiting is broken