

# Operating System: Files and Directories

---

Sang Ho Choi ([shchoi@kw.ac.kr](mailto:shchoi@kw.ac.kr))

School of Computer & Information Engineering  
KwangWoon University

# Persistent Storage

파일이 손상되지 않아도 데이터 유지


- Keep a data **intact** even if there is a power loss
    - Hard disk drive
    - Solid-state storage device
  - Two key abstractions in the virtualization of storage
    - File
    - Directory
- 이 디스크 드라이브에게 보여질 것인가  
암호화된지 → (설명은 뒤에 디스크 암호화?)

# File System

- File system
  - A software for organizing data on a storage device

User's viewpoint

작동 관리



system 파일 관리하는

작동 관리

Storage management's  
viewpoint



HDD



NAND flash memory



SSD


# File System (Cont.)

---

- File system of the user's viewpoint
  - File system interface
  - How to show the file system to user?
  - File, directory, attribute, and operation  
*read, write...*
  - tree structure
  
- File system of the storage management viewpoint
  - File system implementation
  - How to map the logical file system to the storage device?
  - Layout, data structures, and algorithms
  - It is required to understand the storage internals

# File System (Cont.)

- File system's objective




- A linear array of bytes
- Each file has low-level name as **inode number**
  - The user is not aware of this name
- File system has a responsibility to store data persistently on disk

# Directory

- Directory is like a file, also has a low-level name
  - It contains a list of (user-readable name, low-level name) pairs
  - Each entry in a directory refers to either *files* or other *directories*
- Example)
  - A directory has an entry ("foo", "10")
    - A file "foo" with the low-level name "10"

pair

# Directory Tree (Directory Hierarchy)



## An Example Directory Tree

## Valid files (absolute pathname) :

/foo/bar.txt

/bar/foo/bar.txt

## Valid directory :

/  
/food

/ba

/bar/ba

/bar/foo/

## – Sub-directories

# Creating Files

- Use `open()` system call with `O_CREAT` flag

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);
```

- `O_CREAT` : create file
  - `O_WRONLY` : only write to that file while opened
  - `O_TRUNC` : make the file size zero (remove any existing content)
- `open()` system call returns **file descriptor**
- *File descriptor* is an integer, and is used to access files

# Reading and Writing Files

- An Example of reading and writing 'foo' file

```
prompt> echo hello > foo
prompt> cat foo
hello
prompt>
```

- `echo` : redirect the output of echo to the file foo
- `cat` : dump the contents of a file to the screen

How does the cat program access the file foo?

We can use `strace` to trace the system calls made by a program

우리랑 같이 예제를 진행하는 건가?  
`strace` 썩 가능.

# Reading and Writing Files (Cont.)

```
prompt> strace cat foo  
...  
open("foo", O_RDONLY|O_LARGEFILE) = 3  
read(3, "hello\n", 4096)          = 6  
write(1, "hello\n", 6)            = 6 // file descriptor 1: standard out  
hello  
read(3, "", 4096)                = 0 // 0: no bytes left in the file  
close(3)                         = 0  
...  
prompt>
```

*64bit offset*

- open (**file descriptor, flags**)
  - Return file descriptor (3 in example)
  - File descriptor 0, 1, 2, is for standard input/ output/ error
- read (**file descriptor, buffer pointer, the size of the buffer**)
  - Return the number of bytes it read
- write (**file descriptor, buffer pointer, the size of the buffer**)
  - Return the number of bytes it write

# Reading and Writing Files (Cont.)

---

- Writing a file (A similar set of read steps)
  - A file is opened for writing (`open()`)
  - The `write()` system call is called
    - Repeatedly called for larger files
  - `close()`

# Reading And Writing, But Not Sequentially

---

- An open file has a **current offset**
  - Determine where the next read or write will begin reading from or writing to within the file
- Update the current offset
  - **Implicitly:** A read or write of  $N$  bytes takes place,  $N$  is added to the current offset
  - **Explicitly:** `lseek()`

# Reading And Writing, But Not Sequentially (Cont.)

```
off_t lseek(int fildes, off_t offset, int whence);
```

- fildes : File descriptor
- offset : Position the file offset to a particular location within the file
- whence : Determine how the seek is performed
- examples

```
off_t curpos;
curpos = lseek(fd, 0, SEEK_CUR); // get the current offset

lseek(fd, 0, SEEK_SET);
lseek(fd, 0, SEEK_END);
lseek(fd, -10, SEEK_CUR);
lseek(fd, 100, SEEK_END);
```

# Writing Immediately with `fsync()`

---

- The file system will buffer writes in memory for some time
  - Ex) 5 seconds, or 30
  - Performance reasons
- At that later point in time, the write(s) will actually be issued to the storage device
  - Write seem to complete quickly
  - Data can be lost (e.g., the machine crashes)

# Writing Immediately with fsync() (Cont.)

---

- However, some applications require more than eventual guarantee
  - Ex) DBMS requires force writes to disk from time to time  
*Database Management System*
- off\_t fsync(int fd)
  - Filesystem forces all dirty (i.e., not yet written) data to disk for the file referred to by the file description
  - fsync() returns once all of these writes are complete

# Writing Immediately with fsync() (Cont.)

- Example of `fsync()`

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);
assert (fd > -1)
int rc = write(fd, buffer, size);
assert (rc == size);
rc = fsync(fd);
assert (rc == 0);
```

- In some cases, this code needs to `fsync()` the directory that contains the file `foo`

# Renaming Files

- `rename(char* old, char *new)`
  - Rename a file to different name
  - It implemented as an **atomic call**
    - Ex) Change from `foo` to `bar`:

```
prompt> mv foo bar           // mv uses the system call rename()
```

➤ Ex) How to update a file atomically:

```
int fint fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC);
write(fd, buffer, size); // write out new version of file
fsync(fd);
close(fd);
rename("foo.txt.tmp", "foo.txt");
```

# Getting Information About Files

- `stat()`, `fstat()`: Show the file metadata
  - Metadata is information about each file
  - Ex) Size, Low-level name, Permission, ...
  - `stat` structure is below:

```
struct stat {  
    dev_t st_dev;          /* ID of device containing file */  
    ino_t st_ino;          /* inode number */  
    mode_t st_mode;        /* protection */  
    nlink_t st_nlink;      /* number of hard links */  
    uid_t st_uid;          /* user ID of owner */  
    gid_t st_gid;          /* group ID of owner */  
    dev_t st_rdev;         /* device ID (if special file) */  
    off_t st_size;          /* total size, in bytes */  
    blksize_t st_blksize;   /* blocksize for filesystem I/O */  
    blkcnt_t st_blocks;     /* number of blocks allocated */  
    time_t st_atime;        /* time of last access */  
    time_t st_mtime;        /* time of last modification */  
    time_t st_ctime;        /* time of last status change */  
};
```

# Getting Information About Files (Cont.)

- To see stat information, you can use the command line tool stat

```
prompt> echo hello > file
prompt> stat file

File: 'file'
Size: 6 Blocks: 8 IO Block: 4096 regular file
Device: 811h/2065d Inode: 67158084 Links: 1
Access: (0640/-rw-r-----)Uid: (30686/ root) Gid: (30686/ remzi)
Access: 2011-05-03 15:50:20.157594748 -0500
Modify: 2011-05-03 15:50:20.157594748 -0500
Change: 2011-05-03 15:50:20.157594748 -0500
```

- File system keeps this type of information in a inode structure

# Removing Files

- `rm` is Linux command to remove a file
  - `rm` call `unlink()` to remove a file

```
prompt> strace rm foo
...
unlink("foo")          = 0      // return 0 upon success
...
prompt>
```

Why it calls `unlink()`? not “remove or delete”

We can get the answer later

# Making Directories

- `mkdir()`: Make a directory

```
prompt> strace mkdir foo  
...  
mkdir("foo", 0777) = 0  
prompt>
```

- When a directory is created, it is **empty**
- Empty directory have two entries:
  - . (itself), .. (parent)

```
prompt> ls -a  
./      ../  
prompt> ls -al  
total 8  
drwxr-x--- 2 remzi remzi   6 Apr 30 16:17 ./  
drwxr-x--- 26 remzi remzi 4096 Apr 30 16:17 ../
```

# Reading Directories

- A sample code to read directory entries (like ls)

```
int main(int argc, char *argv[]) {
    DIR *dp = opendir(".");
                                // open current directory
    assert(dp != NULL);
    struct dirent *d;
    while ((d = readdir(dp)) != NULL) // read one directory entry
    {
        // print out the name and inode number of each file
        printf("%d %s\n", (int) d->d_ino, d->d_name);
    }
    closedir(dp);
                                // close current directory
    return 0;
}
```

- The information available within struct dirent

```
struct dirent {
    char          d_name[256];      /* filename */
    ino_t         d_ino;           /* inode number */
    off_t         d_off;           /* offset to the next direct */
    unsigned short d_reclen;       /* length of this record */
    unsigned char d_type;          /* type of file */
}
```

# Deleting Directories

---

- `rmdir()`: Delete a directory
    - Require that the directory be empty
    - If you call `rmdir()` to a non-empty directory, it will fail
- I.e., Only has “.” and “..” entries

# Hard Links

- link(old pathname, new one)
  - Link a new file name to an old one
  - Create another way to refer to *the same file*
  - The command-line link program : ln

```
prompt> echo hello > file
prompt> cat file
hello
prompt> ln file file2 // create a hard link, link file to file2
prompt> cat file2
hello
```

# Hard Links (Cont.)

---

- The way link works:
  - Create another name in the directory
  - Refer it to the same inode number of the original file
    - The file is not copied in any way
  - Then, we now just have two human names (`file` and `file2`) that both refer to the same file

# Hard Links (Cont.)

- The result of `link()`

```
prompt> ls -i file file2
67158084 file /* inode value is 67158084 */
67158084 file2 /* inode value is 67158084 */
prompt>
```

- Two files have **same inode number**, but two human name (file, file2)
- There is **no difference** between file and file2
  - Both just links to the underlying metadata about the file

# Hard Links (Cont.)

- Thus, to remove a file, we call `unlink()`

```
prompt> rm file  
removed 'file'  
prompt> cat file2           // Still access the file  
hello
```

- *reference count*

- Track how many different file names have been linked to this inode
  - When `unlink()` is called, the reference count decrements
  - If the reference count reaches zero, the file system free the inode and related data blocks → truly “delete” the file

# Hard Links (Cont.)

- The result of `unlink()`
  - `stat()` shows the reference count of a file

```
prompt> echo hello > file          /* create file*/
prompt> stat file
... Inode: 67158084 Links: 1 ...
prompt> ln file file2              /* hard link file2 */
prompt> stat file
... Inode: 67158084 Links: 2 ...
prompt> stat file2
... Inode: 67158084 Links: 2 ...
prompt> ln file2 file3            /* hard link file3 */
prompt> stat file
... Inode: 67158084 Links: 3 ...
prompt> rm file
/* remove file */
prompt> stat file2
... Inode: 67158084 Links: 2 ...
prompt> rm file2
/* remove file2 */
prompt> stat file3
... Inode: 67158084 Links: 1 ...
prompt> rm file3
/* Link count is 1 */
```

# Symbolic Links (Soft Link)

- Symbolic link is more **useful** than Hard link
  - Hard Link cannot create to a directory
  - Hard Link cannot create to a file to other partition
    - Because inode numbers are only unique within a file system
- Create a symbolic link: `ln -s`

```
prompt> echo hello > file
prompt> ln -s file file2 /* option -s : create a symbolic link, */
prompt> cat file2
hello
```

# Symbolic Links (Cont.)

- What is different between *Symbolic link* and *Hard Link*?
  - Symbolic links are a third type the file system knows about

```
prompt> stat file
... regular file ...
prompt> stat file2
... symbolic link ...           // Actually a file it self of a different type
```

- The size of symbolic link (file2) is 4 bytes

```
prompt> ls -al
drwxr-x---  2 remzi remzi  29 May  3 19:10 .
drwxr-x--- 27 remzi remzi 4096 May  3 15:14 ../          // directory
-rw-r-----  1 remzi remzi   6 May  3 19:10 file           // regular file
1rwxrwxrwx  1 remzi remzi   4 May  3 19:10 file2 -> file    // symbolic link
```

➤ A symbolic link holds the pathname of the linked-to file as the data of the link file

# Symbolic Links (Cont.)

- If we link to a longer pathname, our link file would be bigger

```
prompt> echo hello > alongerfilename
prompt> ln -s alongerfilename file3
prompt> ls -al alongerfilename file3
-rw-r----- 1 remzi remzi 6 May 3 19:17 alongerfilename
lrwxrwxrwx 1 remzi remzi 15 May 3 19:17 file3 -> alongerfilename
```

# Symbolic Links (Cont.)

- Dangling reference
  - When remove an original file, symbolic link points noting

```
prompt> echo hello > file
prompt> ln -s file file2
prompt> cat file2
hello
prompt> rm file          // remove the original file
prompt> cat file2
cat: file2: No such file or directory
```


# Making and Mounting a File System

---

- `mkfs` tool : Make a file system
  - Write an empty file system, starting with *a root directory*, on to a disk partition
  - Input:
    - A device (such as a disk partition, e.g., `/dev/sda1`)
    - A file system type (e.g., `ext3`)
  - Example
    - `prompt> mkfs -t ext3 /dev/sda1`

# Making and Mounting a File System (Cont.)

- `mount()`
  - Take an existing directory as a target **mount point**
  - `prompt> mount -t ext3 /dev/sda1 /usr`



# Making and Mounting a File System (Cont.)

- mount program: show what is mounted on a system

```
/dev/sda1 on / type ext3 (rw)
proc on /proc type proc (rw)
sysfs on /sys type sysfs (rw)
/dev/sda5 on /tmp type ext3 (rw)
/dev/sda7 on /var/vice/cache type ext3 (rw)
tmpfs on /dev/shm type tmpfs (rw)
AFS on /afs type afs (rw)
```

- ext3: A standard disk-based file system
- proc: A file system for accessing information about current processes
- tmpfs: A file system just for temporary files
- AFS: A distributed file system