

Operating System: File System Consistency

Sang Ho Choi (shchoi@kw.ac.kr)

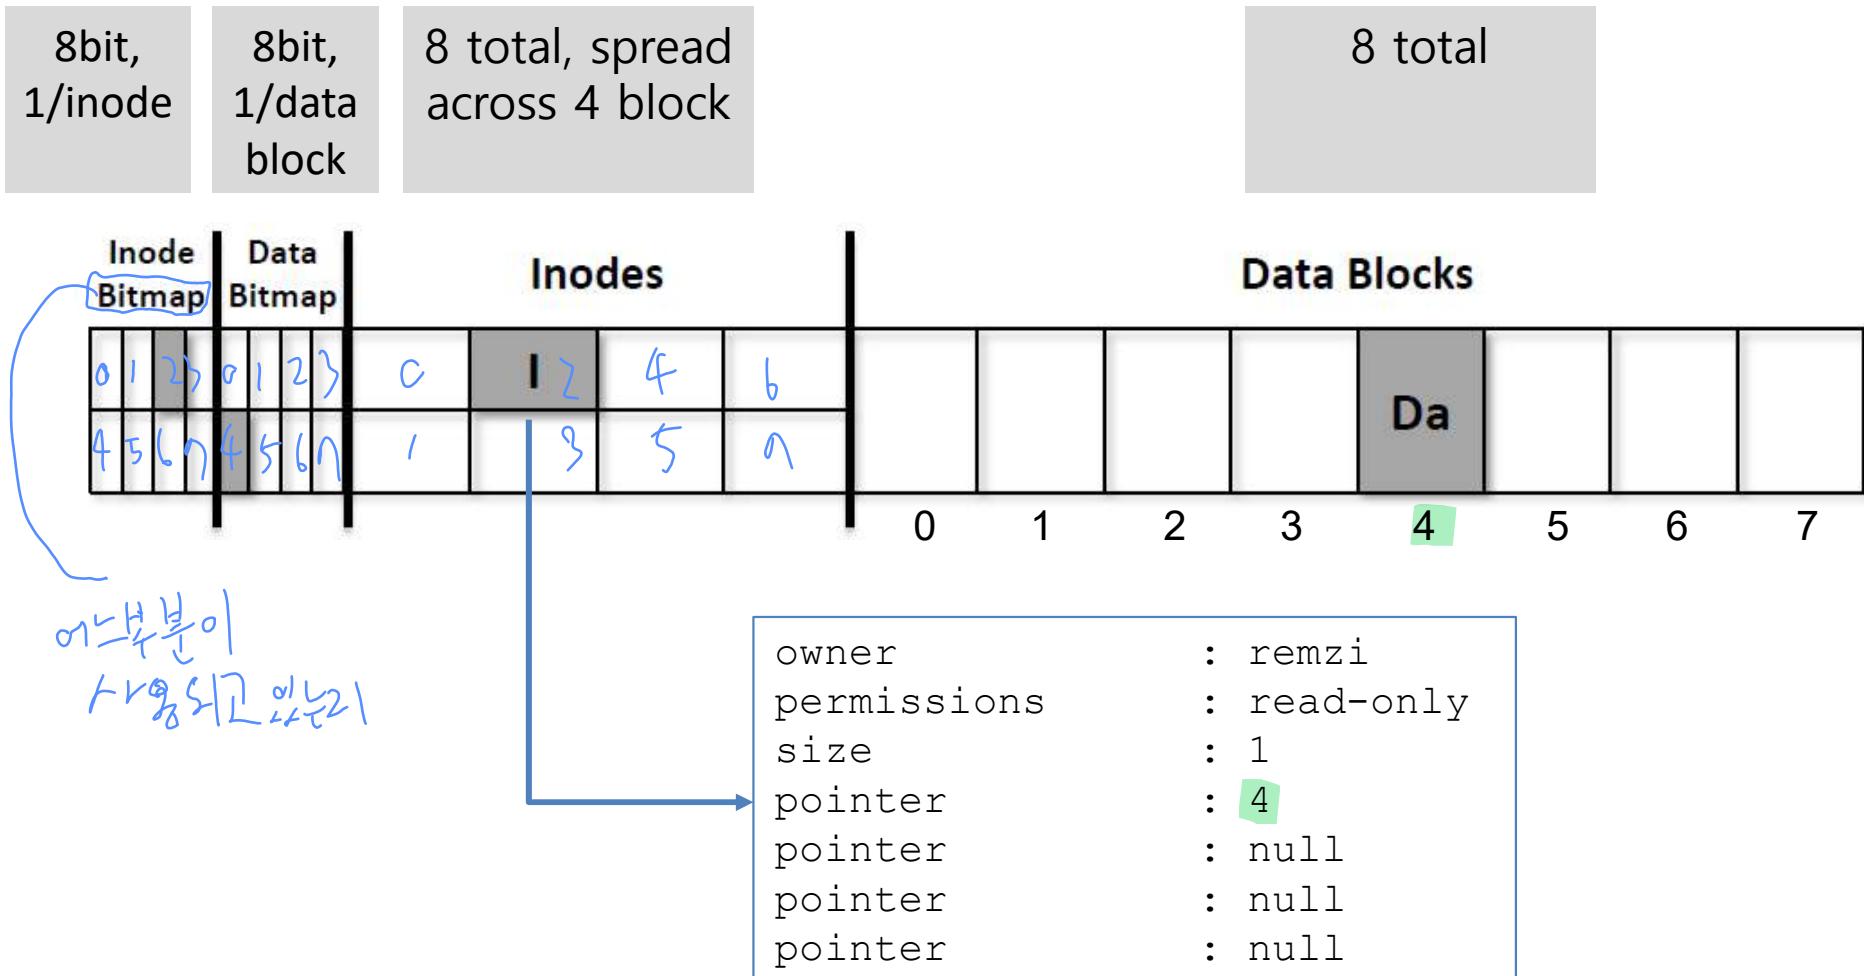
School of Computer & Information Engineering
KwangWoon University

Crash Consistency

- File system may perform several disk writes to complete a single system call
 - e.g. creat(), write(), unlink(), rename(), ...
 - But, disk only guarantees atomicity of a single sector write
- If file system is interrupted between writes, the ondisk structure may be left in an inconsistent state
 - Power loss
 - System crash (kernel panic)
 - Transient hardware malfunctioning
- We want to move file system from one consistent state to another atomically

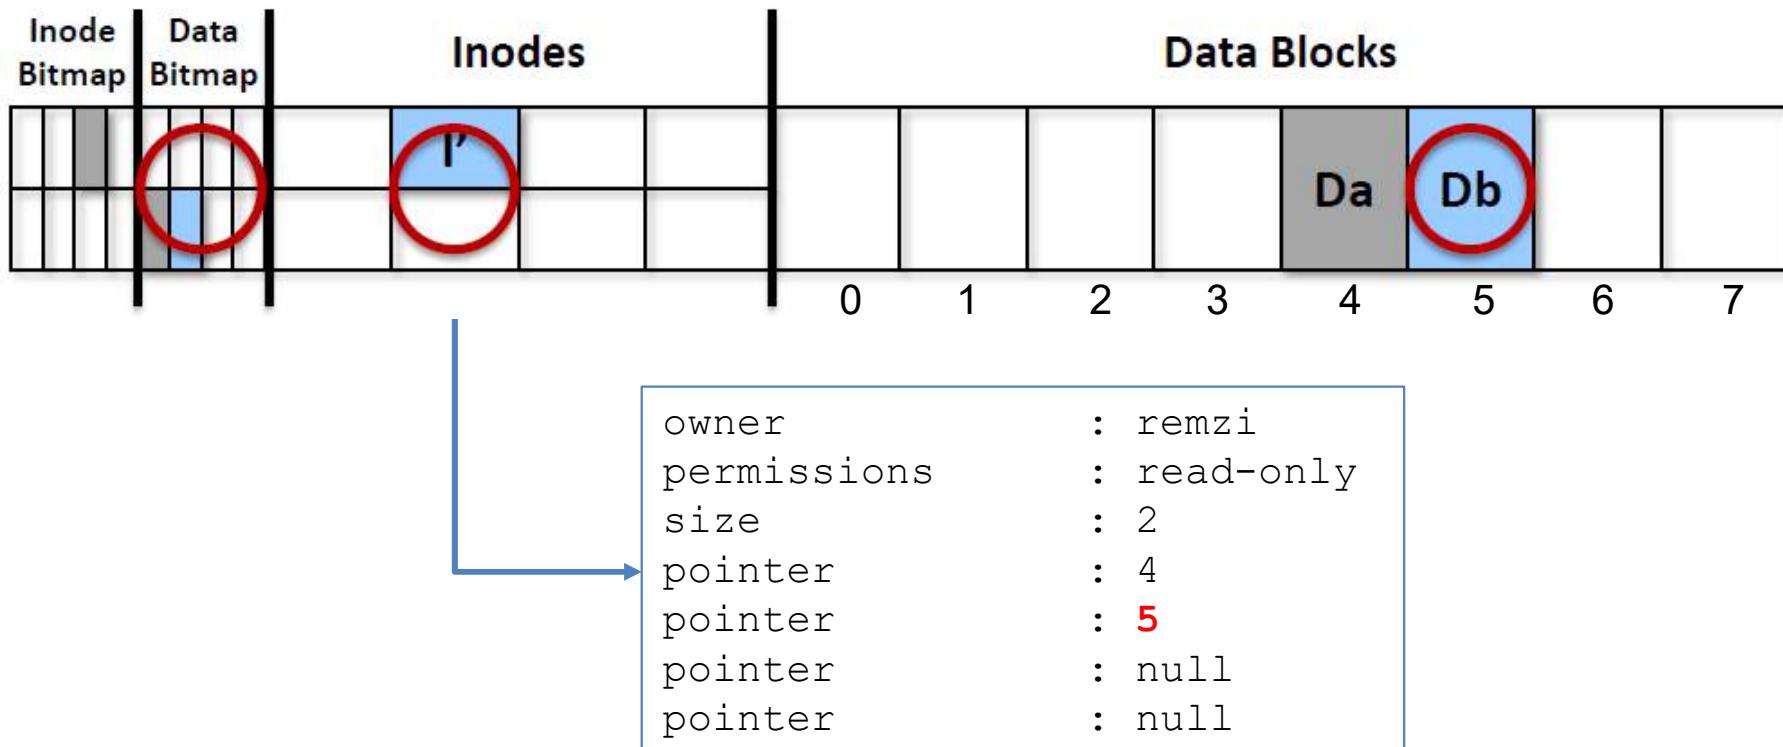
Example: Appending Data

- Initial state



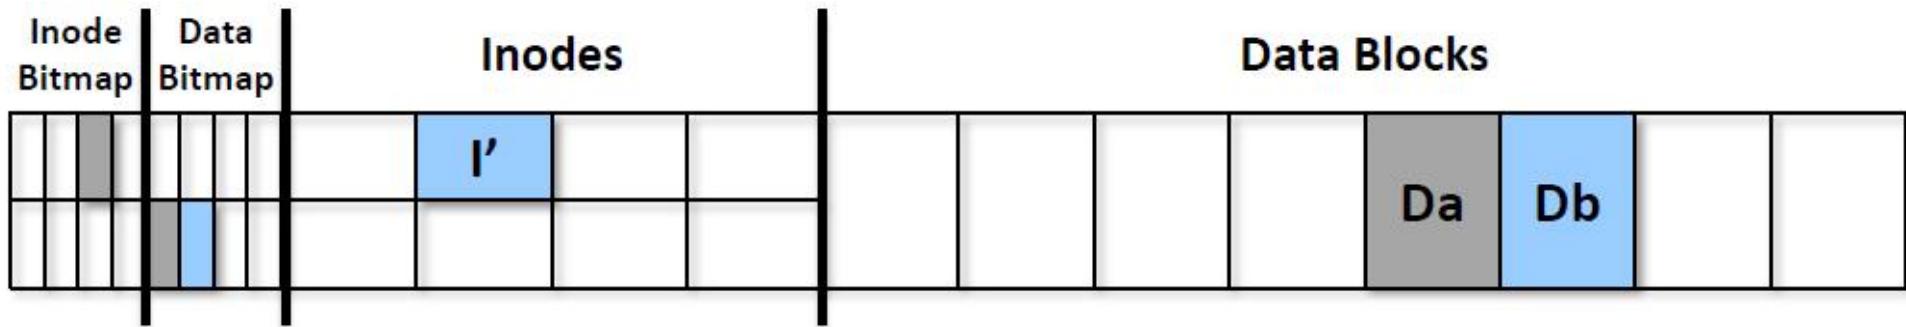
Example: Appending Data

- Appending a data block Db
 - Data bitmap is updated
 - Inode is updated (I')
 - New data block is allocated (Db)

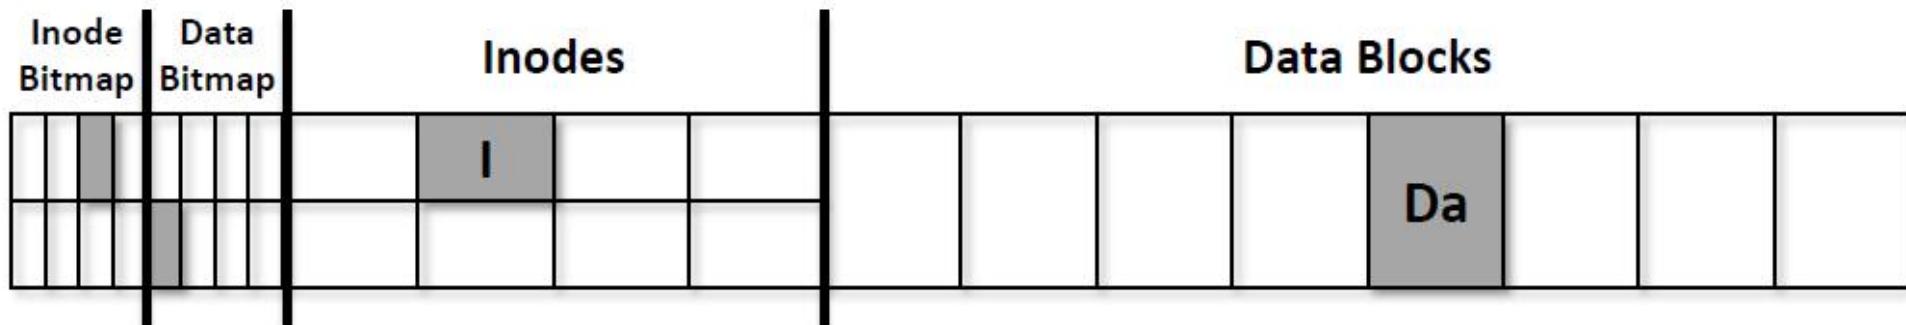


Example: Crash Scenarios (1)

- Everything touched media: **No problem**

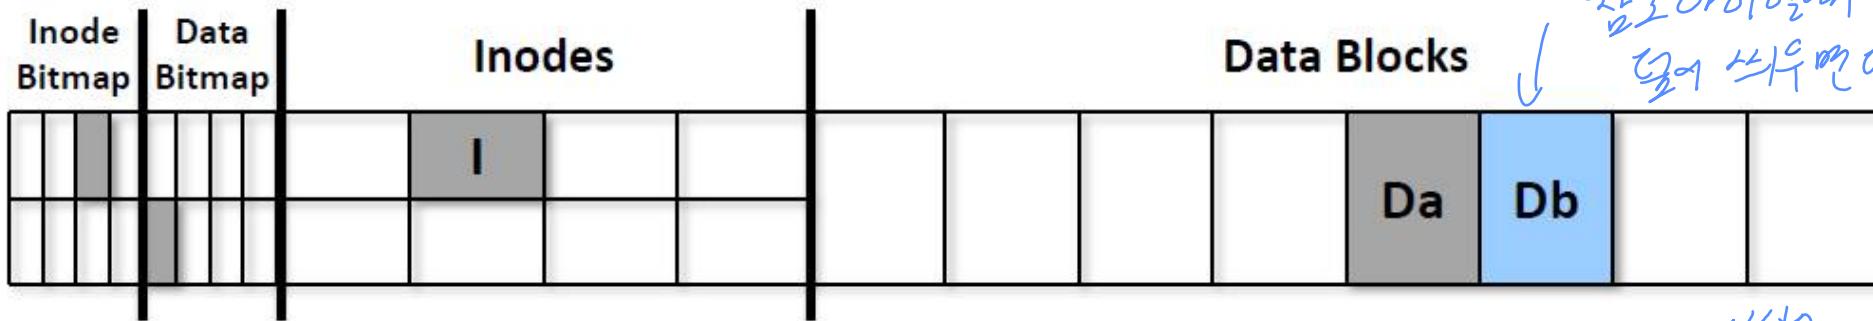


- Nothing touched media: **No problem**
 - Due to page cache or internal disk write buffer



Example: Crash Scenarios (2)

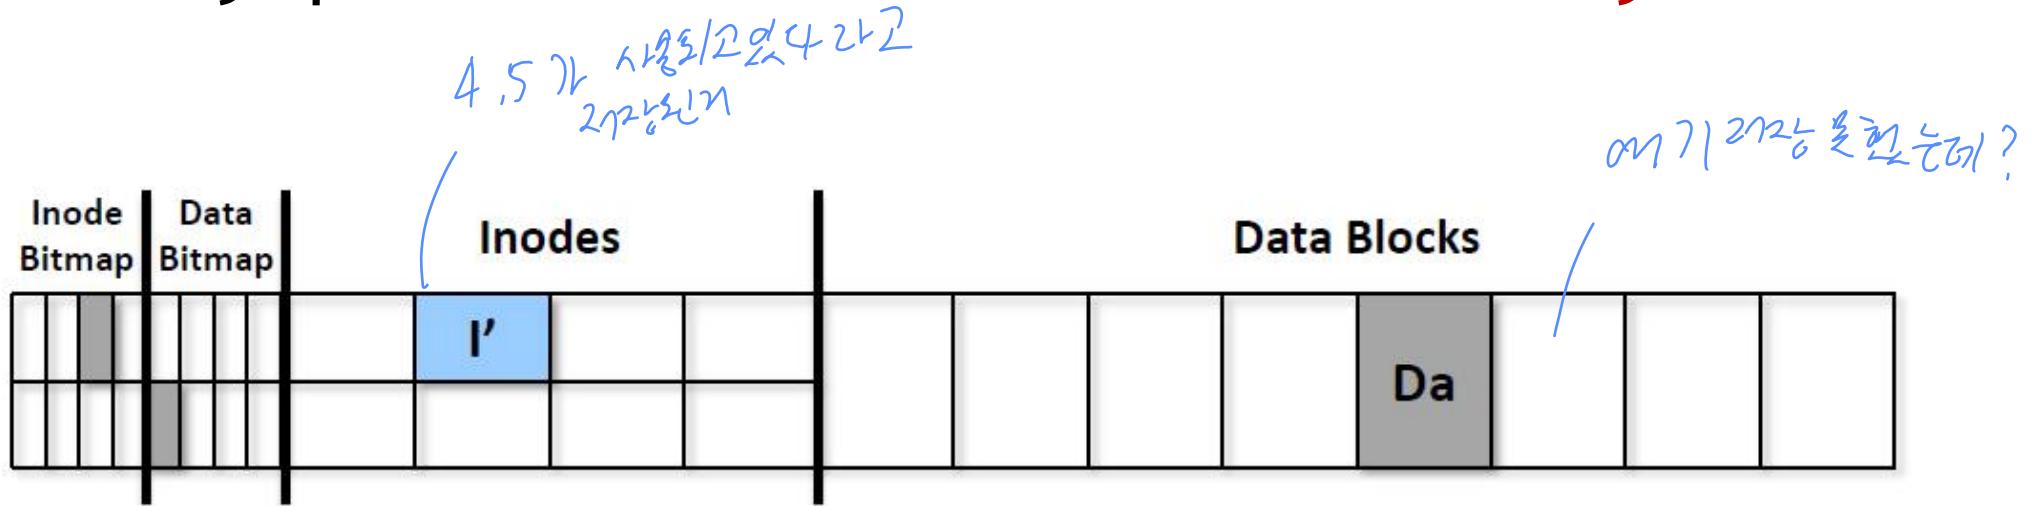
- Only data block (Db) is written: OK



- No inode points to data block 5 (Db)
- Data bitmap says data block 5 is free

Example: Crash Scenarios (3)

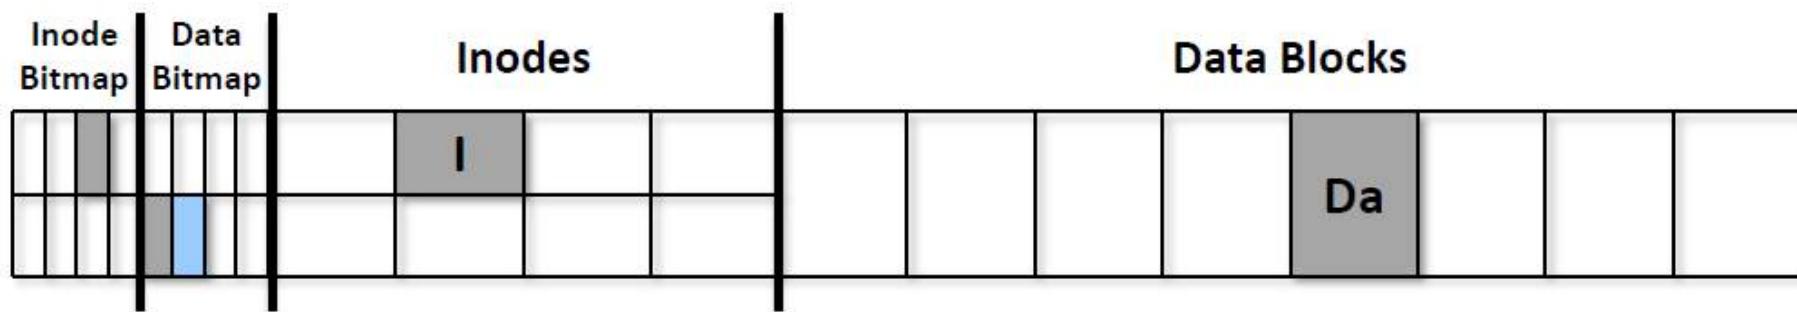
- Only updated inode (I') is written: **Inconsistency**



- Inode I' points to data block 5, but data bitmap says it's free
 - Read will get garbage data (old contents of data block 5)
 - If data block 5 is allocated to another file later, the same block will be used by two inodes

Example: Crash Scenarios (4)

- Only updated data bitmap is written: **Inconsistency**

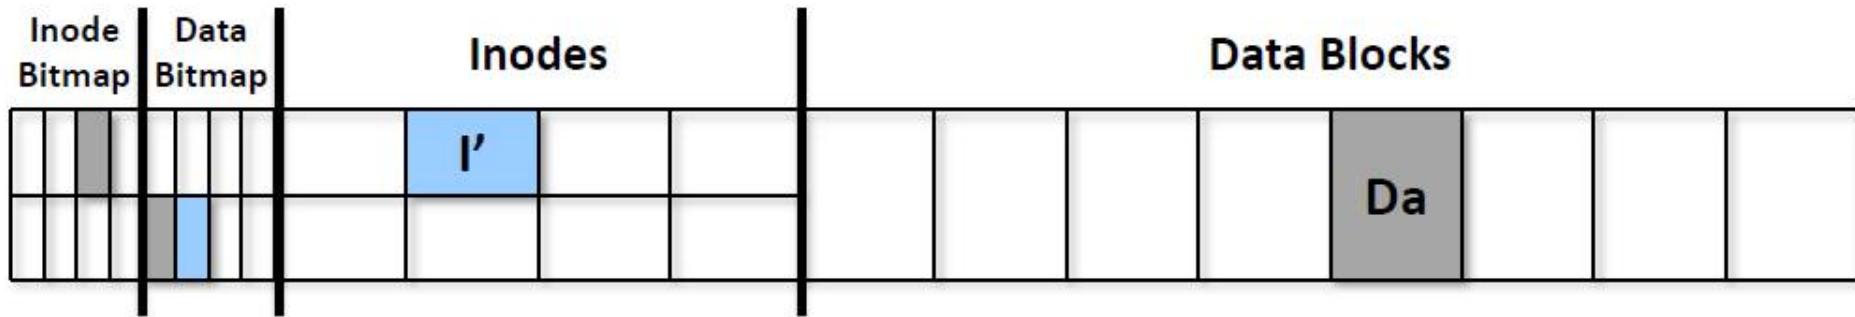


- Data bitmap indicates data block 5 is allocated, but no inode points to it
- Data block 5 will never be used by the file system
- Lost data block (space leak)

Bitmap이 있는 데서도 파일이
존재하지 않는 경우
가장 많은 경우

Example: Crash Scenarios (5)

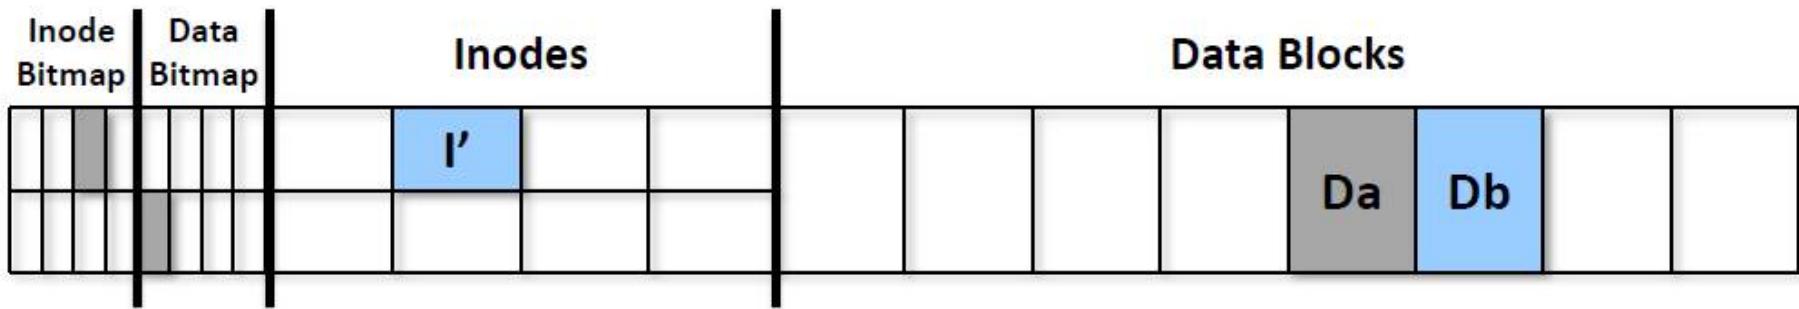
- Only inode and bitmap are written: **OK**



- File system metadata is completely consistent
- Inode I' has a pointer to data block 5 and data bitmap indicates it is in use
- Read will get garbage data (old contents of data block 5)

Example: Crash Scenarios (6)

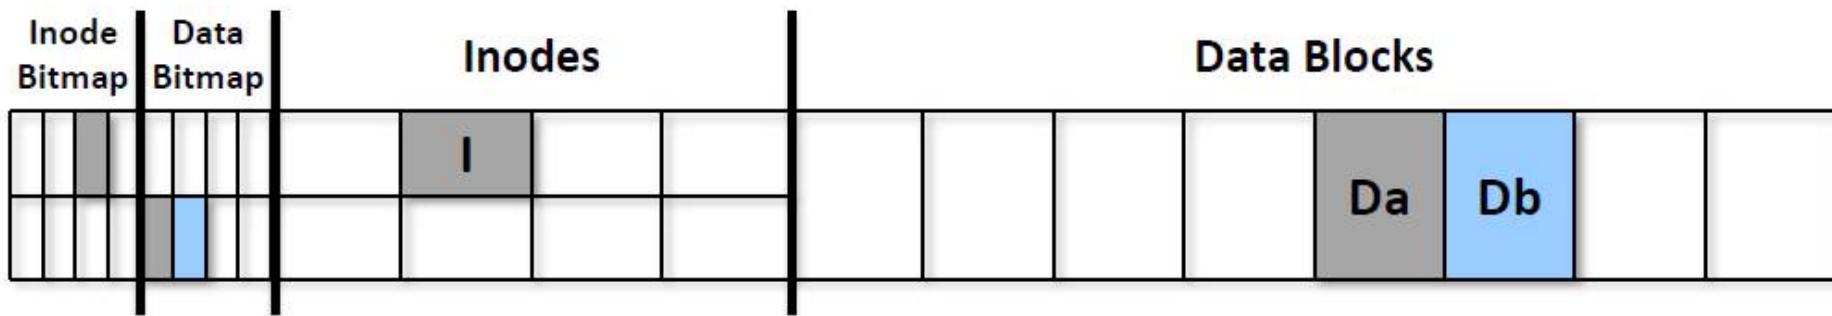
- Only inode and data block are written: **Inconsistency**



- Inode **I'** has a pointer to data block 5, but data bitmap indicates it is free
- Data block 5 can be reallocated to another inode

Example: Crash Scenarios (7)

- Only bitmap and data block are written: **Inconsistency**



- Data bitmap indicates data block 5 is in use, but no inode points to it
- Data block 5 will never be used by the file system
- Lost data block (space leak)

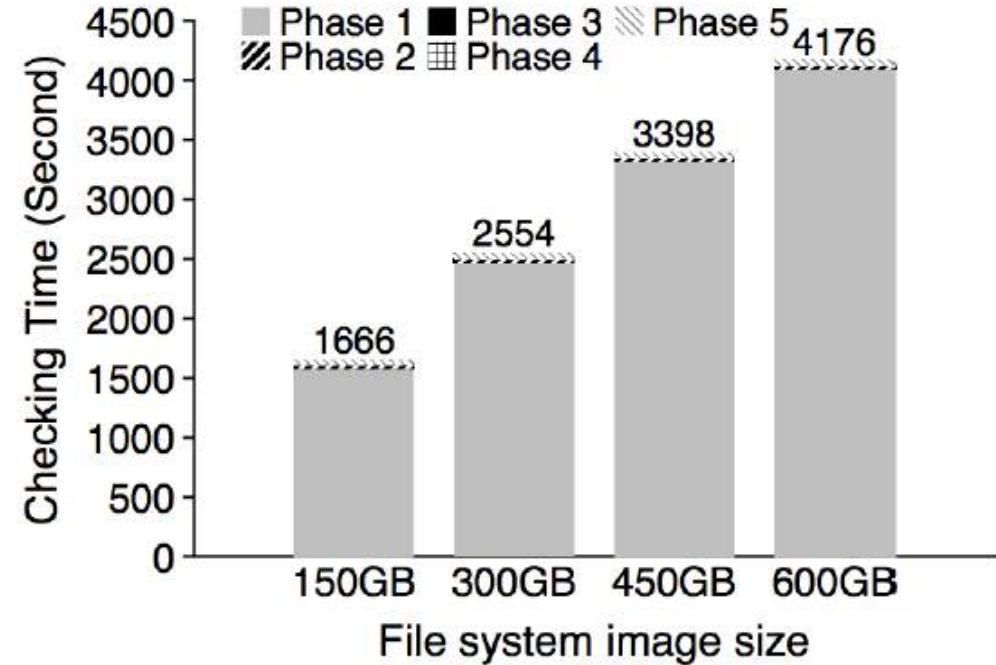
결과적으로 Data bitmap과 inode간 불일치가 없다면 시스템 관점에서의 consistency는 유지가 되는거다.

- File System Checker
 - A Unix tool for finding inconsistencies in a file system and repairing them (cf. Scandisk in Windows) 불일치를 찾아서 고침
 - Run before the file system is mounted and made available
- After crash, scan whole file system for contradictions and “fix” it if needed
 - Inode bitmap consistency
 - Data bitmap consistency
 - Inode link count
 - Duplicated/invalid data block pointers
 - Other integrity checks for superblock, inode, and directories

FSCK Problems

- **Too slow!**

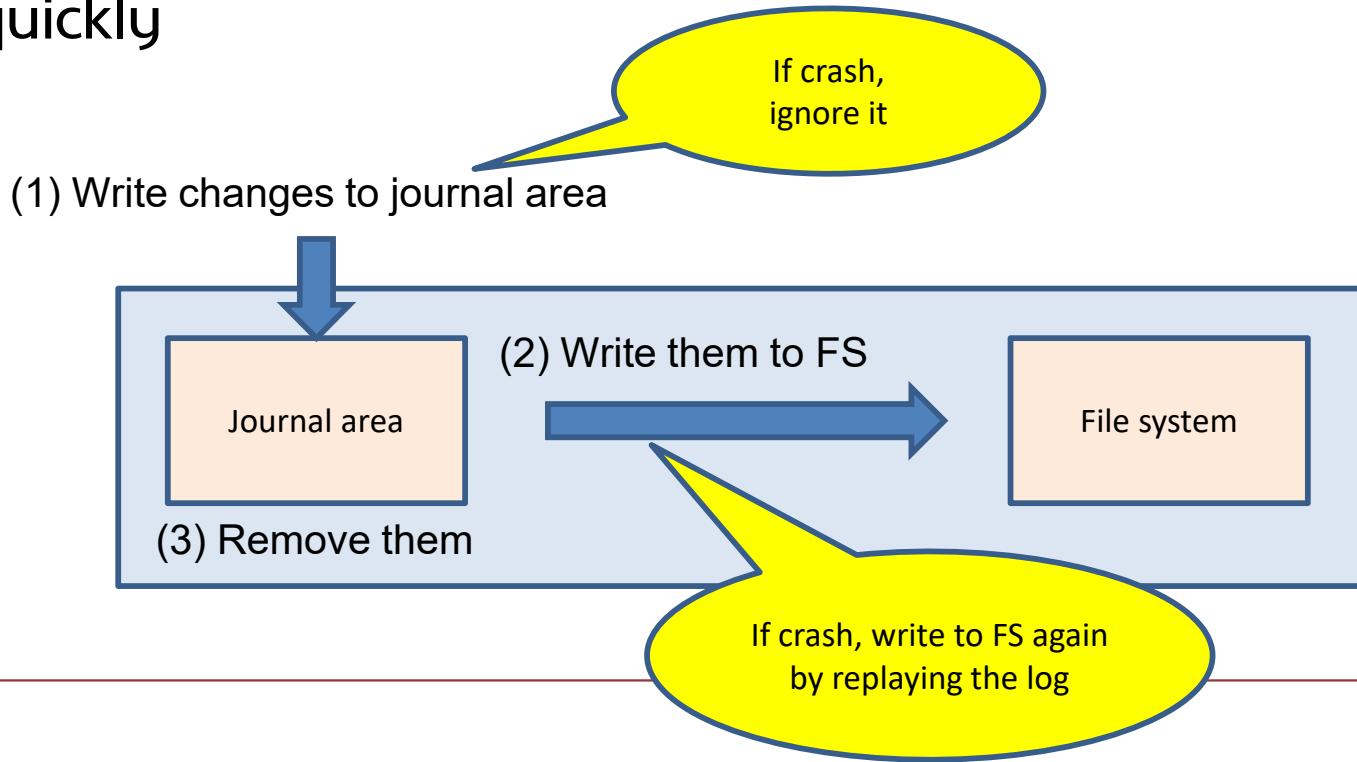
- Need to scan the entire directory tree and block pointers
- Fsck'ing a 600GB disk takes ~ 70 min



- Requires intricate knowledge of the file system
- Not always obvious how to fix file system image
- Don't know "correct" state, just consistent one

Journaling

- Journaling file system
 - If write a "/a/b"?
 - ➔ b's data write, b's inode write, a's data write, a's inode write, bitmap write(if needed), etc.
 - keeps track of changes in a "journal", which is usually a circular log
 - When a system crash, file systems can be recovered more quickly



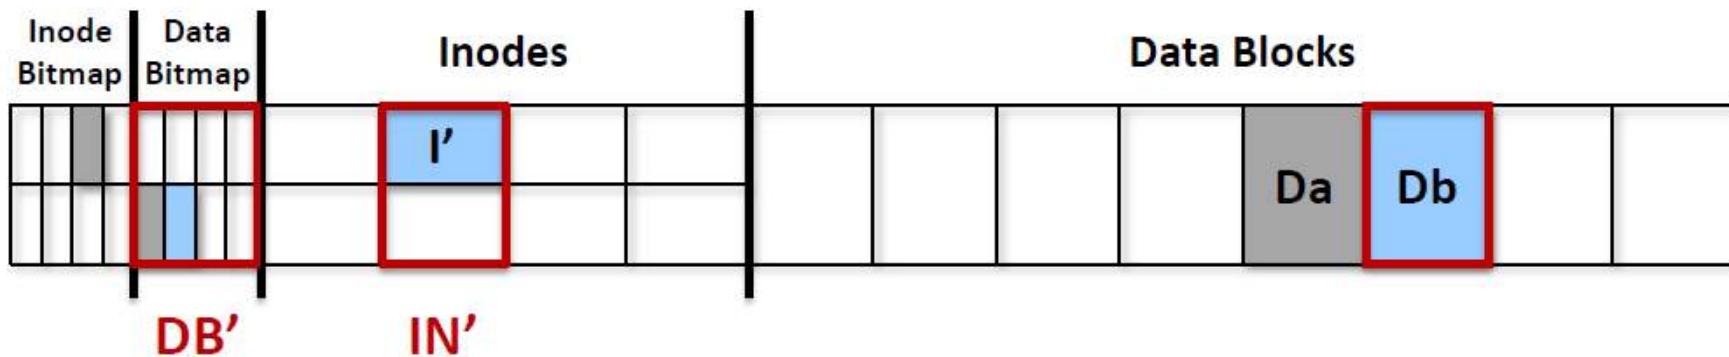
Journaling (Cont.)

- Write-ahead logging
 - A well-known technique for database transactions
 - Record a log, or journal, of changes made to on-disk data structures to a separate location ("journaling area")
 - Write updates to their final locations ("checkpointing") only after the journal is safely written to disk
 - If a crash occurs:
 - Discard the journal if the journal write is not committed
 - Otherwise, redo the updates based on the journal data
 - Fast as it requires to scan only the journaling area
 - Used in modern file systems:
 - Linux Ext3/4, ReiserFS, IBM JFS, SGI XFS, Windows NTFS, ...

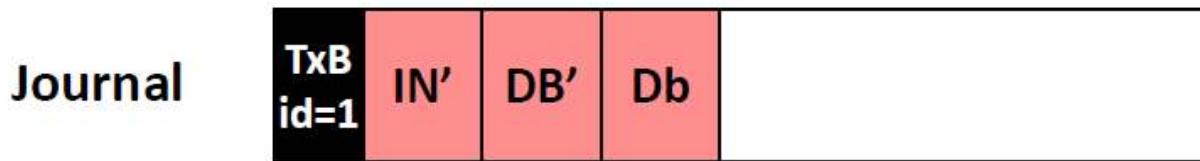
저널에 쓰다 크래쉬면
무시 저널에 있는걸로
쓸라고 할땐 다시쓰기

Example: Appending Data (1)

- Appending a data block Db



- Step 1: Journal write
 - Write journal header block (TxB), inode block (IN'), data bitmap block (DB') and data block (Db)



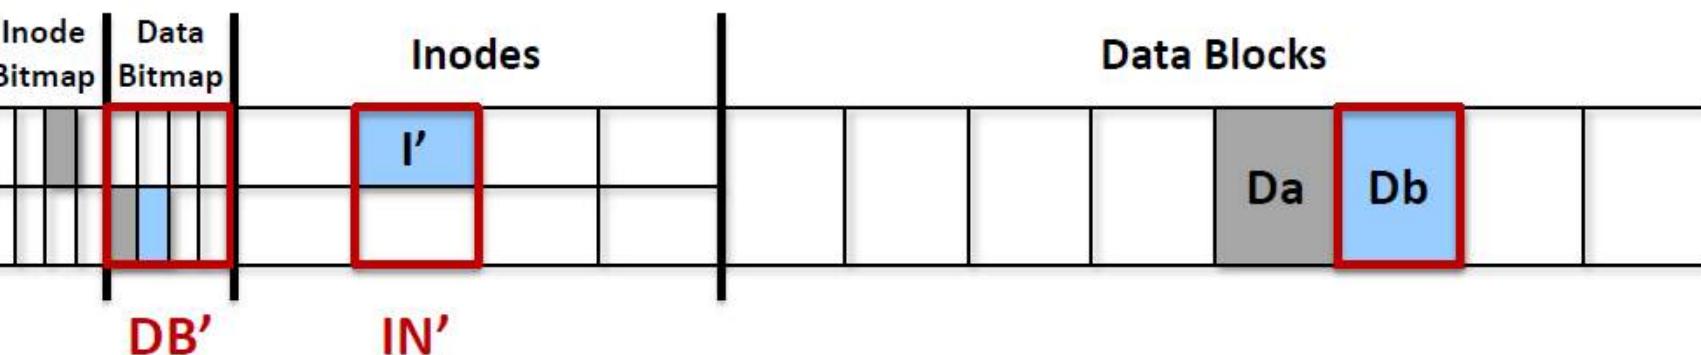
Example: Appending Data (2)

- Step 2: Journal commit
 - Write journal commit block (TxE)



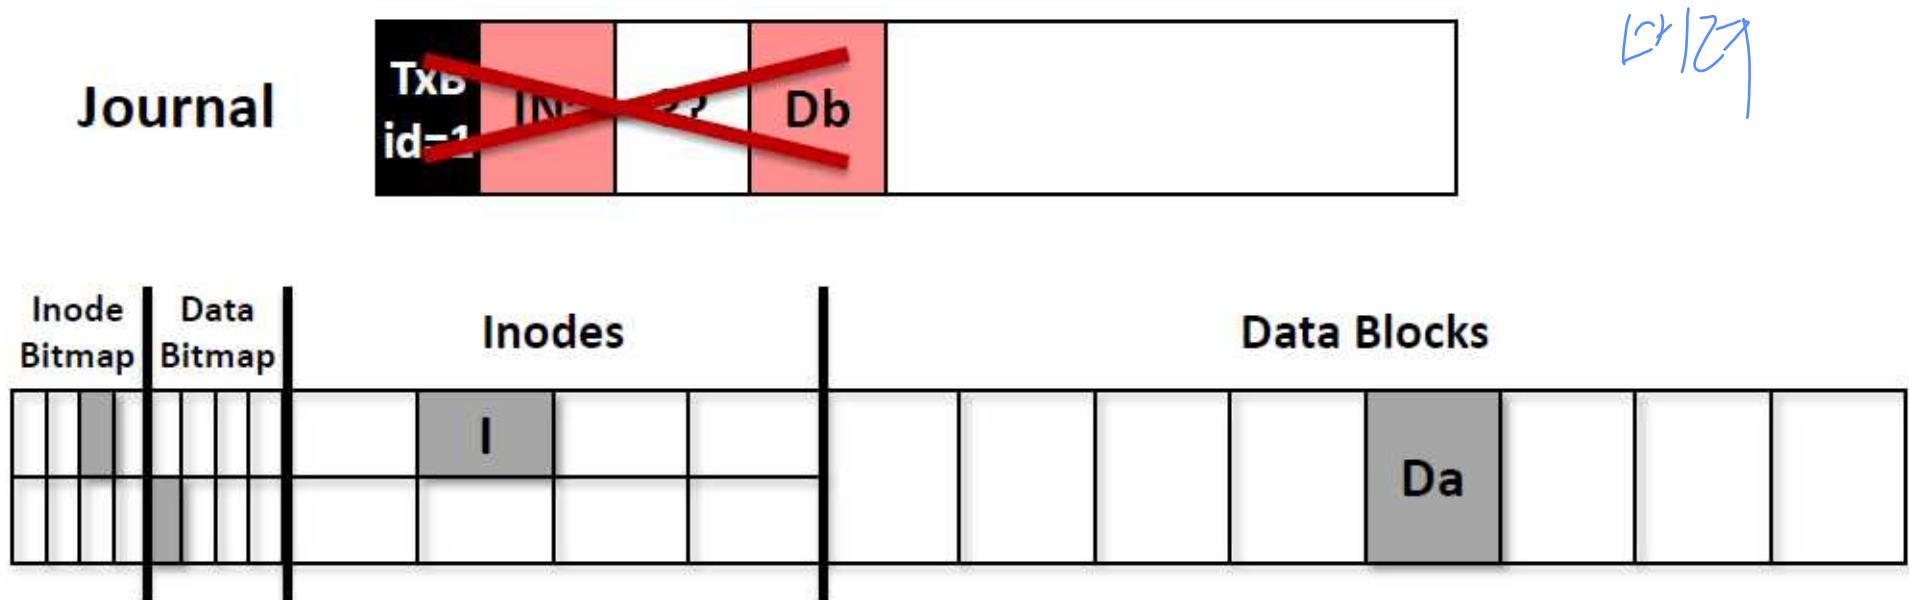
- Step 3: Checkpoint

↑
작업을 완료했습니다.



Example: Recovery (1)

- Crash between step 1 & 2
 - Journal write has not been committed
 - Simply discard the journal
 - File system is rolled back to the state before data block Db is appended



Example: Recovery (2)

CJW/KH

- Crash between step 2 & 3
 - Doesn't matter which metadata/data blocks were actually updated
 - Roll-forward recovery (redo logging): overwrite their final on-disk locations using the journal data

